


Chapter 4: Threads





Chapter 4: Threads

- Overview
- Implementation of Thread
- Multithreading Models
- Threading issues
- Operating System Examples





Overview

- To build **parallel** programs, such as:
 - Parallel execution on a **multiprocessor**
 - **Web server** to handle multiple simultaneous web requests
- We will need to:
 - **Create** several processes that can execute in parallel
 - The OS will then **schedule** these processes in parallel





Processes Overheads

- A full process includes numerous things:
 - an address space (defining all the code and data pages)
 - OS resources and accounting information
 - a “thread of control”,
 - ▶ defines where the process is currently executing
 - ▶ That is the PC and registers
- ☞ Creating a new process is **costly**
 - all of the structures (e.g., page tables) that must be allocated
- ☞ **Communicating** between processes is costly
 - most communication goes through the OS





Need “Lightweight” Processes

- What’s similar in these processes?
 - They all share the **same code and data (address space)**
 - They all share the **same privileges**
 - They share almost everything in the process
- What don’t they share?
 - Each has its own PC, registers, and stack pointer
- Idea: why don’t we separate the idea of process (address space, accounting, etc.) from that of the minimal “thread of control” (PC, SP, registers)?





Threads and Processes

- Most operating systems therefore support two entities:
 - the process,
 - ▶ which defines the address space and general process attributes
 - the thread,
 - ▶ which defines a sequential execution stream **within a process**
- Threads are the unit of **scheduling**
- Processes are *containers* in which threads execute





Overview

- A traditional process is basic unit of program concurrent execution
 - It is a basic unit to own resources
 - It is a basic unit of CPU scheduling
- So, a traditional process is called **heavyweight** process
- However, most modern OS separate the two attributes of process
 - Process is a unit to own **resources**
 - Thread is a basic unit of **CPU utilization**
- So , a thread is called a **lightweight** process





Overview

- A *thread* (or *lightweight process*) is a basic unit of CPU utilization
 - program counter
 - register set
 - stack space
- A thread shares with its peer threads its:
 - **code** section
 - **data** section
 - **operating-system resources**, such as open files and signals
 - **collectively known as a task.**





Overview

- Many software packages running on modern PCs are multithreaded
- An application is typically implemented as a separate process with several threads of control.
 - A web server might have one thread **display** images or text while another thread **retrieves** data from the network
 - A word processor may have one thread for **displaying** graphics, another thread for **reading** keystrokes from the user, and a third thread for performing spelling and grammar **checking** in the background.






Overview

- There are also situations where **a single application** may be required to perform **several similar tasks**.
 - E.g. a web server accepts client requests for web pages, images, sound, and so forth.
- If the web server ran as a traditional process, it would be able to service only one client at a time with its single process.
 - A client might **wait a long time** to be serviced
 - It may be more **efficient** for one process that contains several threads to serve the same purpose
- Multithread the web server
 - One thread **listening** for client requests
 - Create another thread to **service** the request when request coming.






Single and Multithreaded Processes



single-threaded process

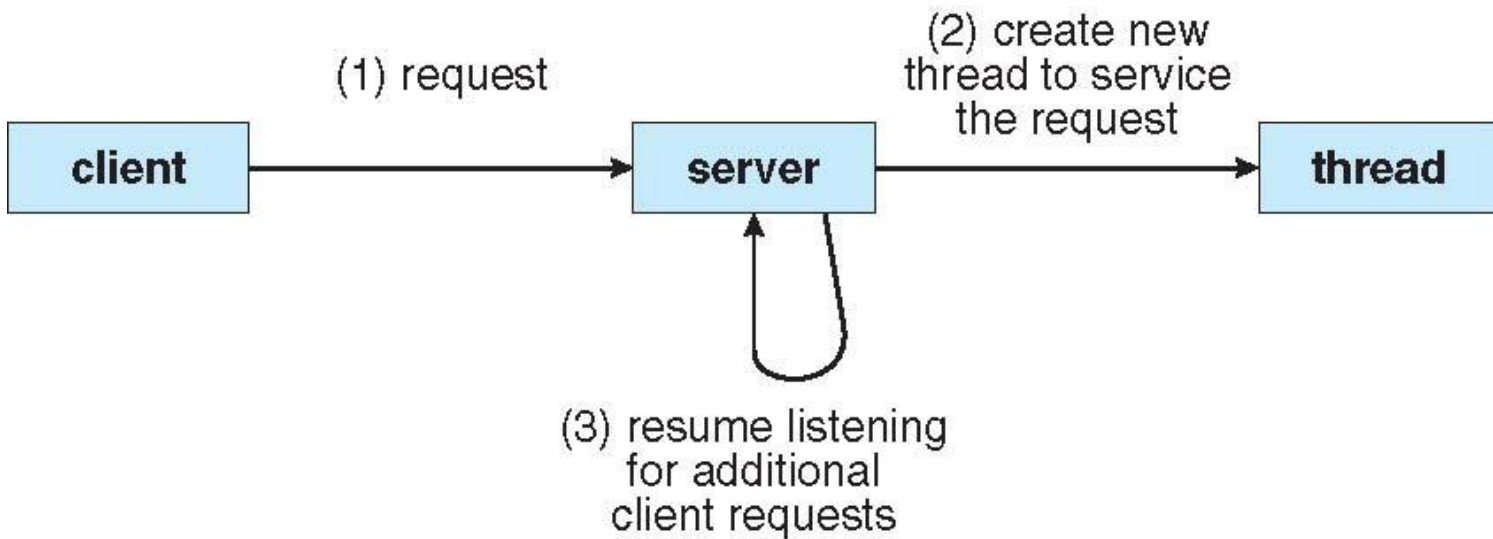


multithreaded process





Multithreaded Server Architecture






Threads vs. Processes

- A thread has **no data segment or heap**
- A thread **cannot live on its own**, it must live within a process
- *There can be more than one thread in a process*
- Inexpensive creation
- Inexpensive context switching
- If a thread dies, its stack is reclaimed
- A process has code/data/heap & other segments
- There **must be at least one** thread in a process
- Threads within a process **share** code/data/heap, share I/O, but each has **its own stack & registers**
- Expensive creation
- Expensive context switching
- If a process dies, its resources are reclaimed & all threads die

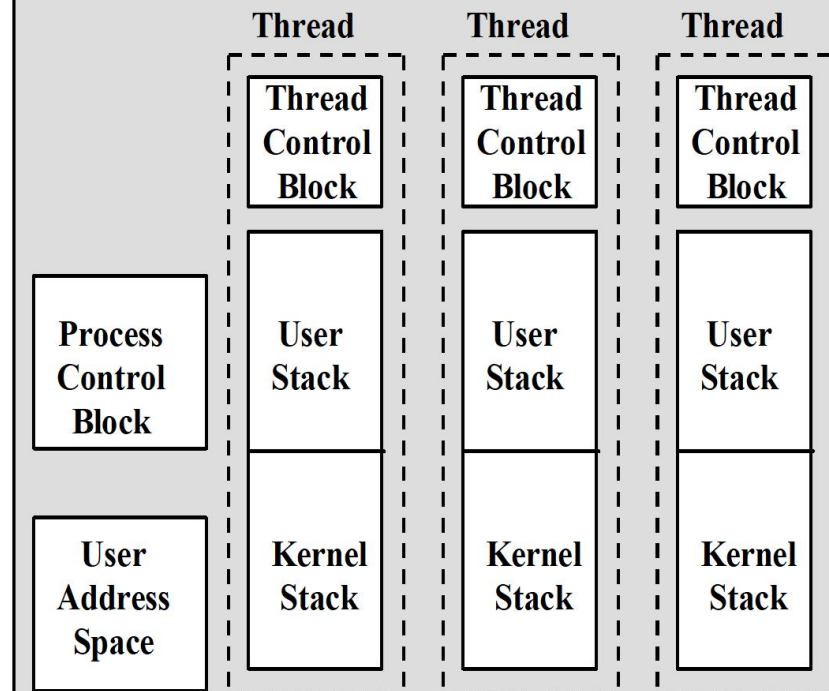




Single-Threaded Process Model



Multithreaded Process Model





Separating Threads and Processes

- Concurrency (multithreading) is useful for:
 - improving program structure
 - handling concurrent events (e.g., web requests)
 - building parallel programs
 - Resource sharing
 - Multiprocessor utilization
- Is multithreading useful even on a single processor?





Benefits

- Responsiveness
- Resource Sharing
- Economy(overhead)
- Scalability





Implementation of thread

- **Kernel-supported** threads
 - Supported directly by the operating system
- **User-level** threads
 - supported above the kernel, via a set of **library calls** at the user level





User-Level Threads


- Thread management done by user-level threads library
 - The library provides support for thread creation, scheduling, and management with no support from the kernel
- Each thread is represented simply by:
 - PC
 - Registers
 - Stack
 - Small control block
- All thread operations are at the user-level:
 - Creating a new thread
 - switching between threads
 - synchronizing between threads





User-Level Threads

- the thread **scheduler is part of a library, outside** the kernel
- thread context switching and scheduling is done by the library
- Can either use cooperative or preemptive threads
 - cooperative threads are implemented by:
 - ▶ CreateThread(), DestroyThread(), Yield(), Suspend(), etc.
 - preemptive threads are implemented **with a timer (signal)**
 - ▶ where the timer handler decides which thread to run next





User-Level Threads

- Three primary thread libraries:
 - POSIX **Pthreads**
 - Win32 threads
 - Java threads





Example User Thread Interface

◆ `t = thread_fork(initial context):`

create a new thread of control;

◆ `thread_stop():`

stop the calling thread, sometimes called `thread_block`;

◆ `thread_start(t):` start the named thread;

◆ `thread_yield():` voluntarily give up the processor;

◆ `thread_exit():`

terminate the calling thread, sometimes called
`thread_destroy`;





Kernel Threads

- Supported by the Kernel
 - Thread is the **unit of CPU scheduling**
- Examples
 - Windows XP/2000
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X





User-Level vs. Kernel Threads

User-Level

- Managed by application
- Kernel not aware of thread
- Context switching cheap
- **Create as many as needed**
- Must be used with care

Kernel-Level

- Managed by kernel
- Consumes kernel resources
- Context switching expensive
- **Number limited** by kernel resources
- Simpler to use

Key issue: kernel threads provide virtual processors to user-level threads, but if all of kthreads block, then all user-level threads will block *even* if the program logic allows them to proceed





Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many
- Two level model





Many-to-One


- Many user-level threads mapped to single kernel thread
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads





Many-to-One Model

Many user-level threads
mapped to single kernel
thread



- ❖ **No parallel** execution of threads - can't exploit multiple CPUs
- ❖ All threads **block** when one uses synchronous I/O






One-to-One

- Each user-level thread maps to kernel thread
- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later





One-to-one Model



- More concurrency
- Better multiprocessor performance
- Each user thread requires creation of kernel thread
- Each thread requires kernel resources; **limits number** of total threads






Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Examples
 - Solaris prior to version 9
 - Windows NT/2000 with the *ThreadFiber* package





Many-to-Many Model



If $U < k$? No benefits of multithreading

If $U > k$, some threads may have to wait for an Kthread to run

- Active thread - executing on an Kthread
- Runnable thread - waiting for an Kthread

A thread **gives up** control of Kthread under the following:

- synchronization, lower priority, yielding, time slicing






Two-level Model

- Similar to M:N($M \geq N$), except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier





Two-level Model



- Supports **both** bound and unbound threads
 - Bound threads - permanently mapped to a single, **dedicated** kthread
 - Unbound threads - may move among kthreads in set
- Thread creation, scheduling, synchronization done in user space





Thread Libraries

- **Thread library** provides programmer with **API** for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS
- Examples:
 - POSIX Pthreads
 - Win32 threads
 - Java thread





Pthreads

- May be provided either as **user-level or kernel-level**
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)





Win32 Threads

- Win32 API is the primary API for Microsoft OS (Win95,98,NT,2000,XP)
- A **kernel-level** library on windows systems






Java Threads

- Java threads are managed by the **JVM**
- Typically implemented using the threads model provided by **underlying OS**
- Java threads may be created by:
 - Extending Thread class
 - Implementing the **Runnable interface**





Java Thread States





Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Thread cancellation of target thread
 - Asynchronous or deferred
- Signal handling
- Thread pools
- Thread-specific data
- Scheduler activations





Semantics of fork() and exec()

- The semantics of fork() and exec() system calls change in a multithreaded program.
- If one thread calls **fork()**, does the new process **duplicate** only the **calling** thread or **all** threads?
- Some UNIX have chosen to have two versions of fork():
 - One duplicates all threads
 - Another duplicates only the thread calling fork()
- The **exec()** typically works in the same way as described in Chapter 3.
 - If a thread calls **exec()**, the program specified in the parameter to exec() will replace the **entire process----** including **all threads**.





Thread Cancellation

- **Terminating** a thread **before** it has finished
 - E.g. a user presses a button on a web browser that stops a web page from loading any further
 - Usually, a web page is loaded using **several** threads----each image is loaded in a separate thread, when the browser is closed, all threads loading the page are canceled.
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread **immediately** ? ?
 - **Deferred cancellation** allows the **target** thread to **periodically check** if it should be cancelled, allowing it an opportunity to terminate itself in an orderly fashion.





Signal Handling

- Signals are used in UNIX systems to notify a process that a particular event has occurred.
- A **signal handler** is used to process signals, follow the same pattern:
 1. Signal is **generated by particular event**
 2. Signal is **delivered to a process**
 3. Signal is handled **once received**





Signal Handling(Cont.)

- Handling signals in a single-threaded program is straightforward
 - Signals are always delivered to a process
- In a multithreaded program, where should a signal be delivered?
 - The following options exist:
 - ▶ Deliver the signal to the thread to **which** the signal **applies**
 - ▶ Deliver the signal to **every** thread in the process
 - ▶ Deliver the signal to **certain** threads in the process
 - ▶ Assign a **specific** thread to **receive all signals** for the process





Thread Pools

- Create a number of threads at process startup and place them into a pool where they await work.
- Advantages:
 - Usually slightly **faster** to service a request with an existing thread than create a new thread
 - **Limits** the number of threads that exist at any one point





Thread Specific Data

- Threads belonging to a process share the data of the process.
- Sometimes, each thread might need its **own copy of certain data**----thread specific data
 - E.g. in a transaction processing system, we might service each transaction in a separate thread, so we need thread specific data
- Most thread libraries (Win32, Pthread, Java) provide some form of support for thread specific data





Scheduler Activations

- Many OS implementing M:M or Two-level models place an **intermediate data structure** between the user and kernel threads-----**lightweight process(LWP)**
- To the user-thread library, the LWP appears to be a **virtual processor** on which the application can schedule a user thread to run
- Each LWP is attached to a kernel thread, and **it is kernel threads that OS schedules to run on physical processors.**
- If a kernel thread **blocks**, the LWP **blocks** as well, and the user-level thread attached to the LWP also **blocks**.





Scheduler Activations

- Scheduler activation is one scheme for **communication** between the user-thread library and the kernel.
- Scheduler activations provide **upcalls** (回调) - a communication mechanism from the kernel to the thread library
- Upcalls are handled by the thread library with an **upcall handler** (回调处理程序), and an upcall handler must run on a virtual processor(LWP).





Operating System Examples

- Windows XP Threads
- Linux Threads
- Solaris threads





Windows XP Threads

- Implements the one-to-one mapping, kernel-level
- Each thread contains
 - A thread id
 - Register set
 - Separate user and kernel stacks
 - Private data storage area
- The register set, stacks, and private storage area are known as the **context** of the threads





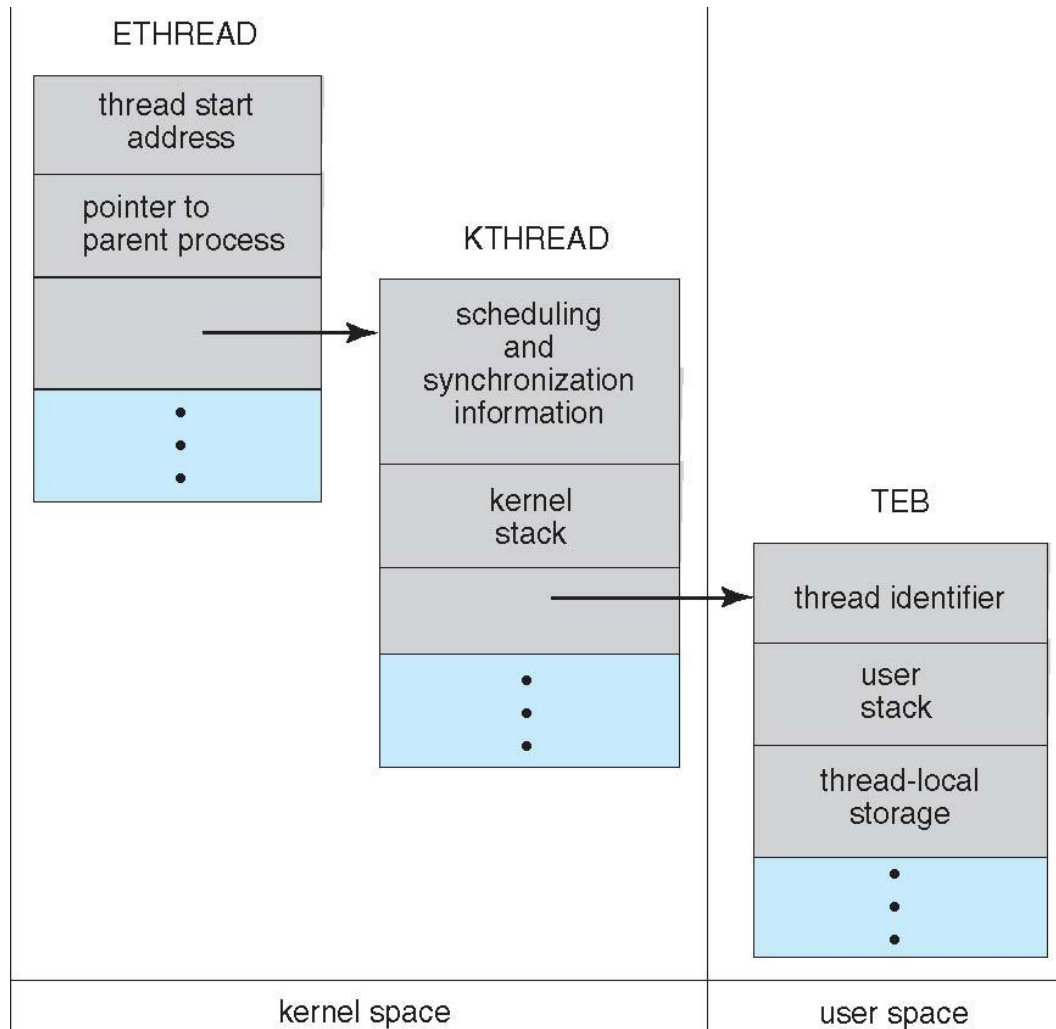
Windows XP Threads

- The primary data structures of a thread include:
 - ETHREAD (executive thread block)
 - KTHREAD (kernel thread block)
 - TEB (thread environment block)





Windows XP Threads





Process Control Block (PCB)

Process ID (PID)
Parent PID
...
Next Process Block •
List of open files •
Image File Name
List of Thread Control Blocks •
...


- This is an abstract view
- Windows implementation of PCB is split in multiple data structures

PCB

Handle Table

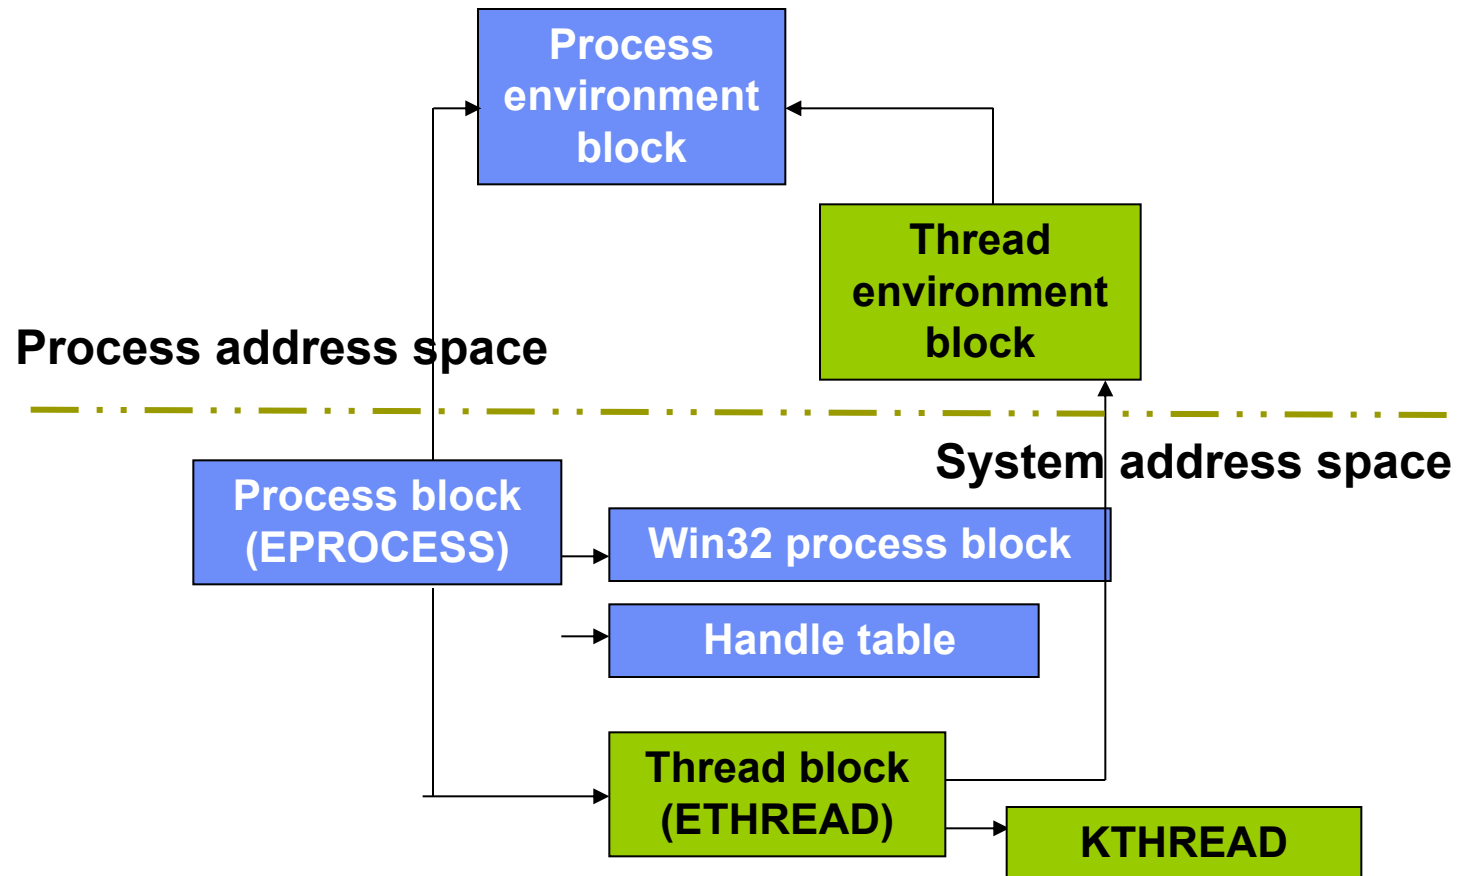
Thread Control Block (TCB)

Next TCB •
Program Counter
Registers
...





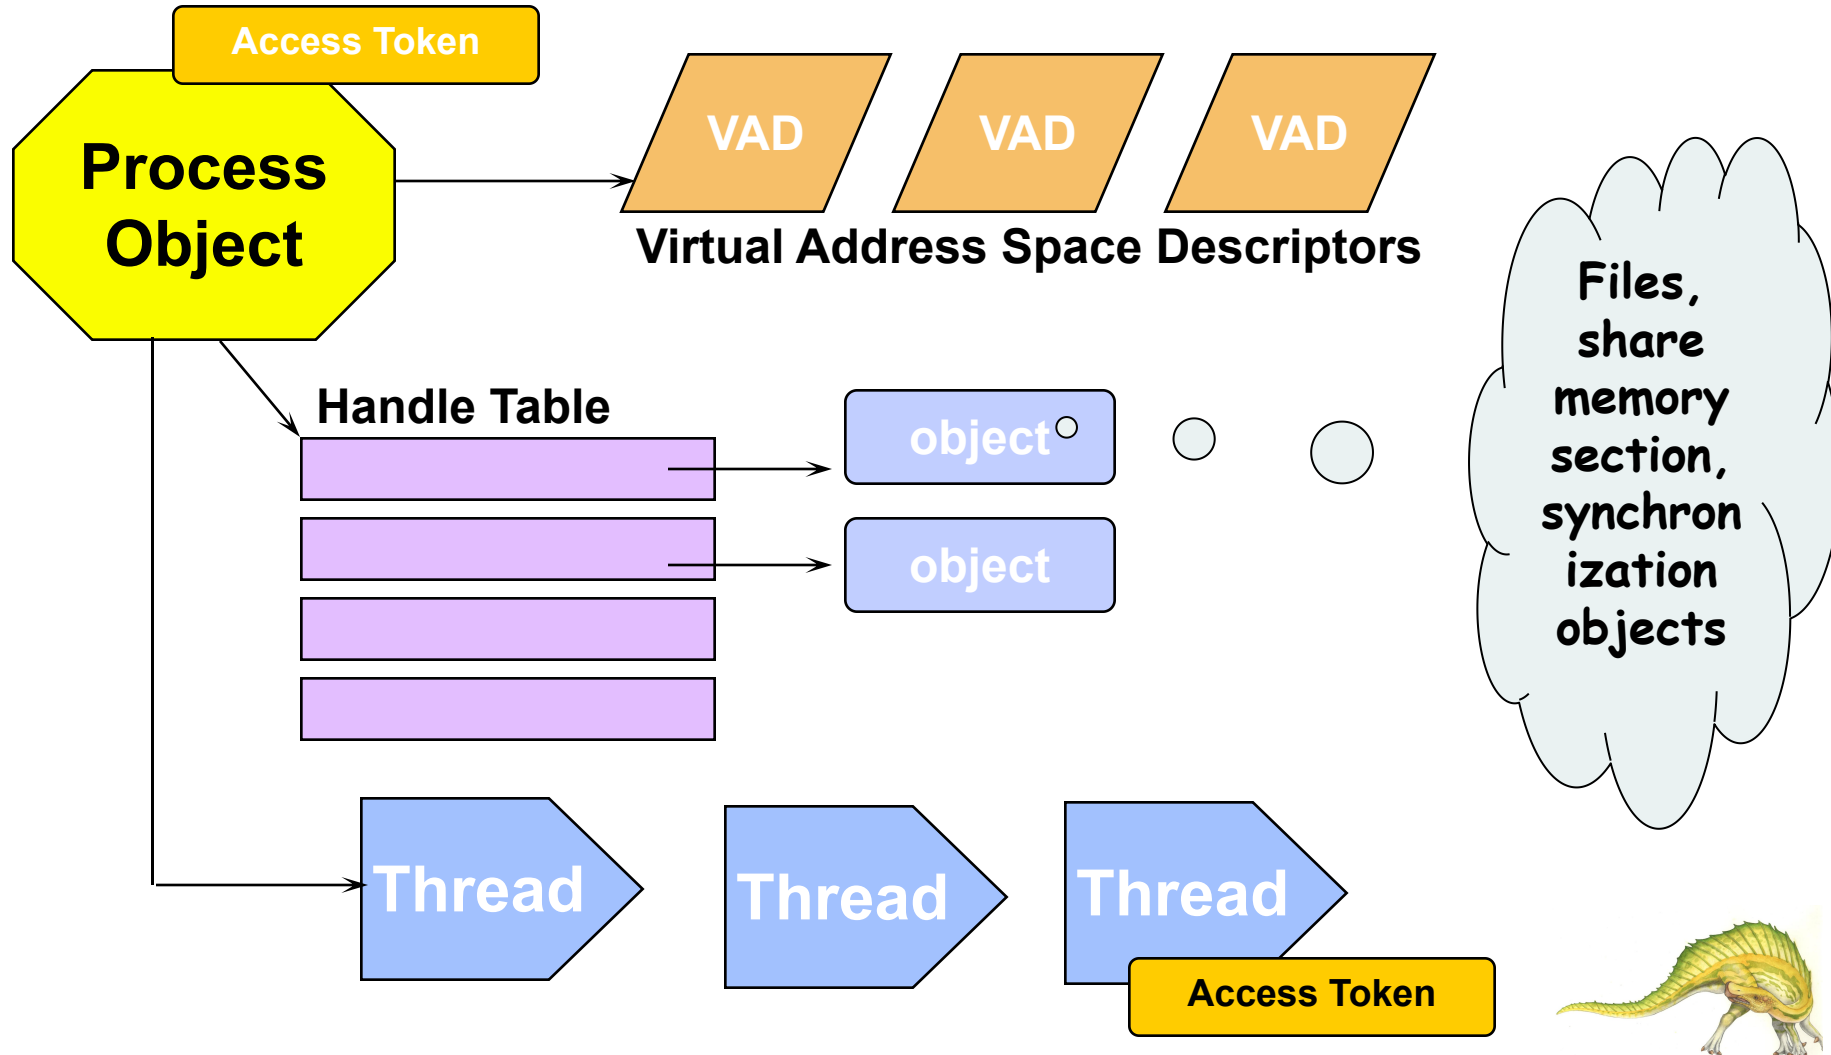
Windows Process and Thread Internals





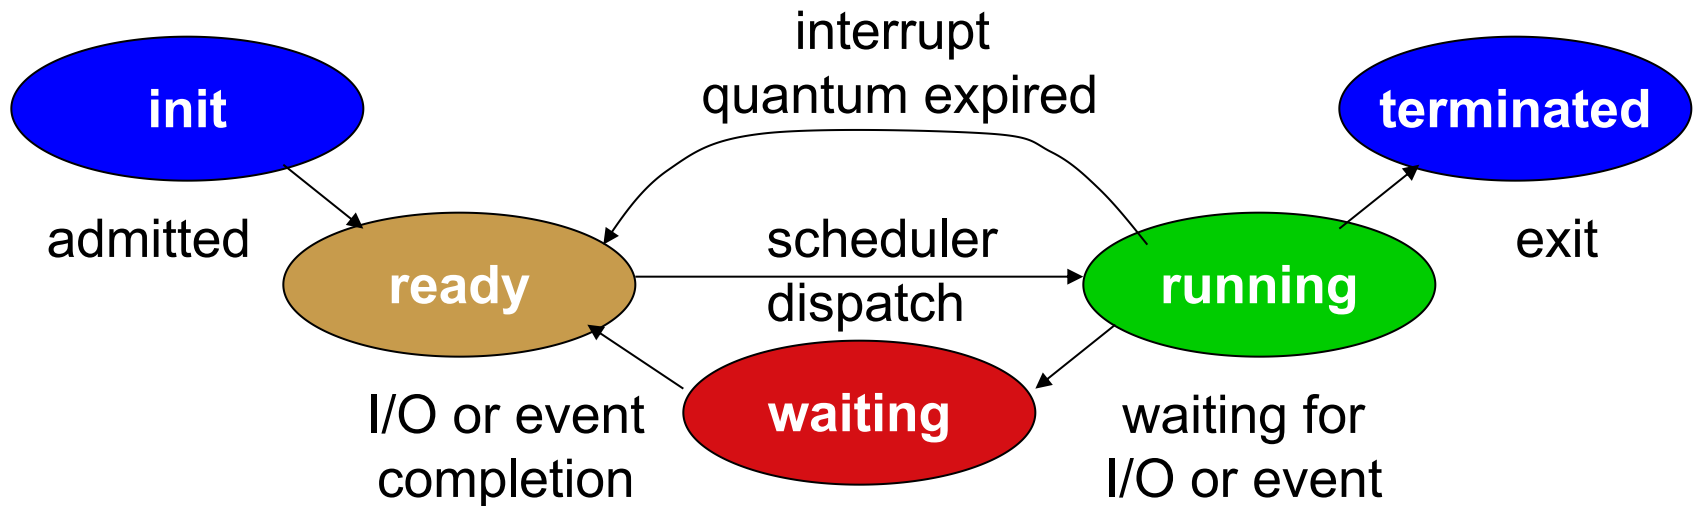
Windows Processes & Threads

Internal Data Structures



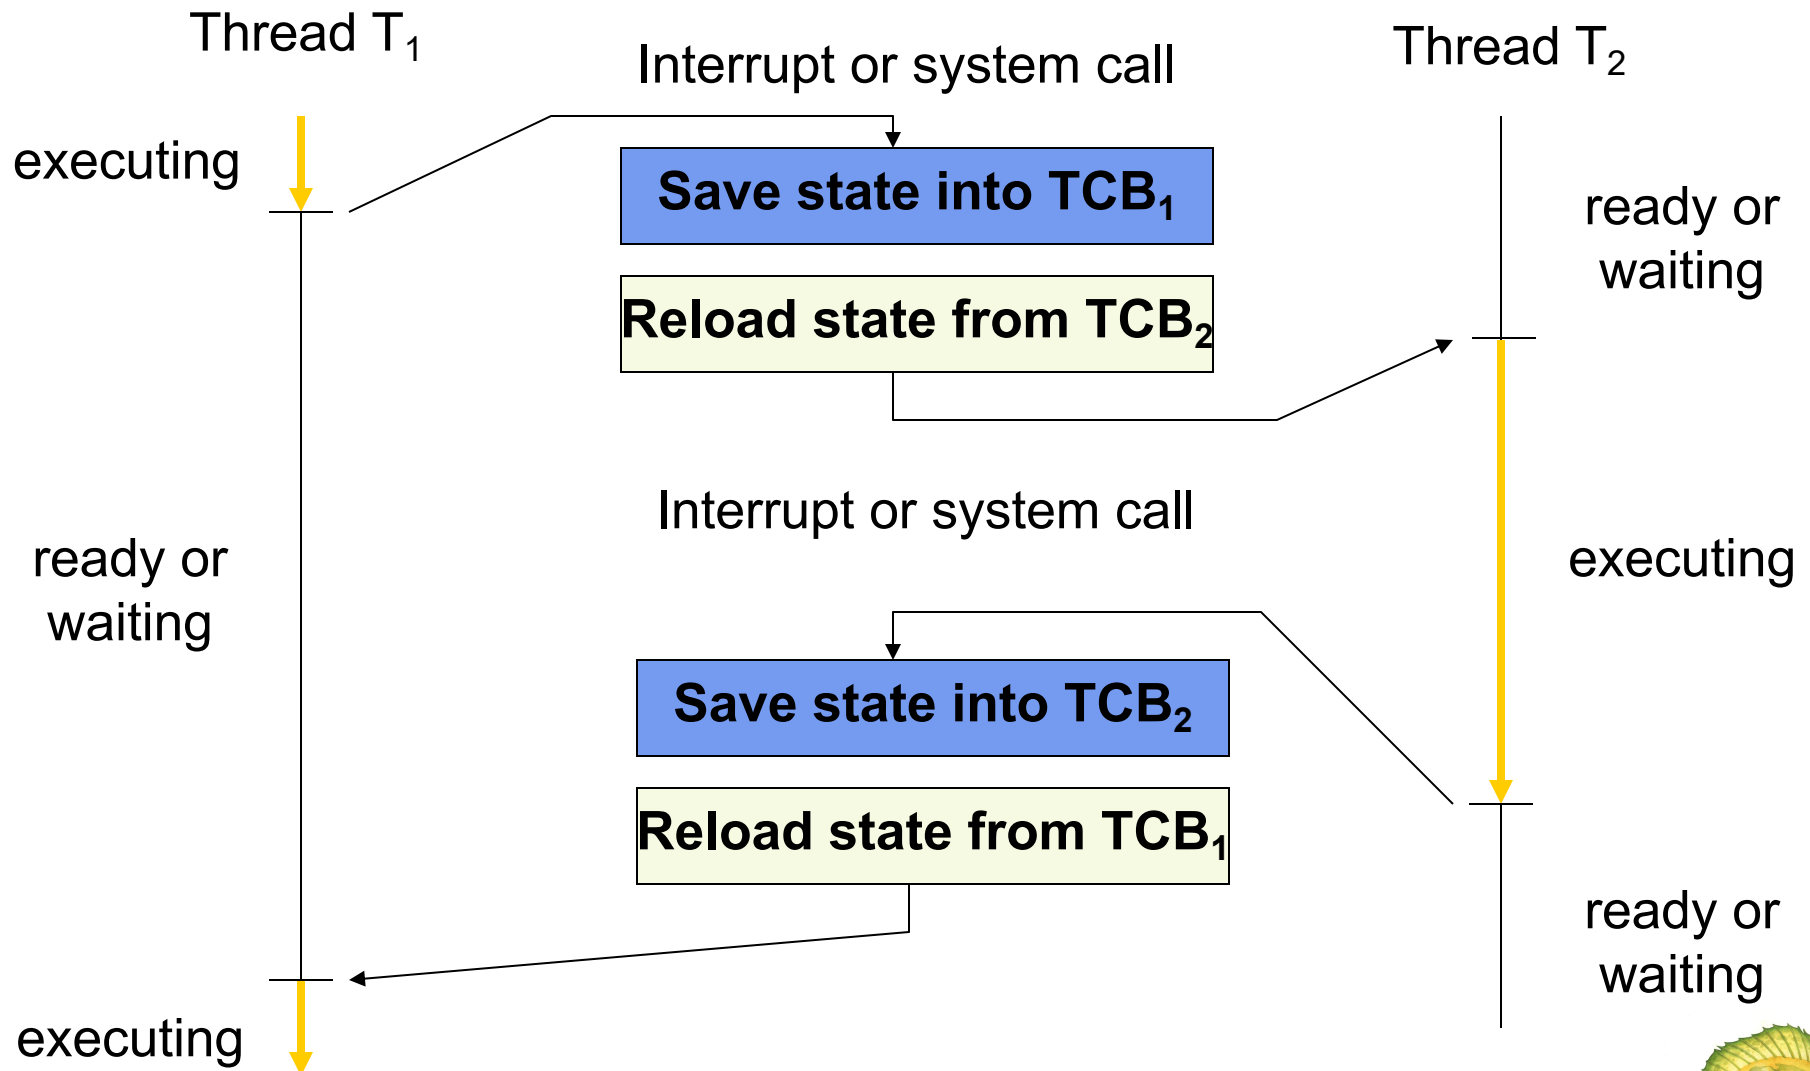


Windows XP Threads





CPU Switch from Thread to Thread





Linux Threads

- Linux does not distinguish between processes and threads
- Linux refers to them as *tasks* rather than *threads*
- Thread creation is done through **clone()** system call
- **clone()** allows a child task to share the address space of the parent task (process)





Linux Threads

- When `fork()`, a task is created, along with a **copy** of all the associated data structures of the parent process
- When `clone()`, a task is created, **no copying** of data structures, the task **points** to the data structures of the parent





Linux Threads

- **clone()** allows a child task to share the address space of the parent task (process)
- **Flags** determine **how much** sharing is to take place between the parent and child
- If no flag is set when clone(), no sharing take place, resulting in functionality similar to fork()

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.





Solaris threads

- Solaris is a version of UNIX with support for threads at the kernel and user levels, symmetric multiprocessing, and real-time scheduling.






Threads Support in Solaris

- Solaris supports
 - Kernel threads
 - Lightweight Processes
 - User Level Threads
- LWP – (lightweight processes) intermediate level between user-level threads and kernel-level threads.





Solaris threads




Solaris process





Solaris threads





Solaris threads

■ Kernel Threads

- The only objects scheduled within the system

■ LWP

- Each process contains at least one LWP
- The thread library multiplexes user-level threads on the pool of LWPs for the process
- Without LWPs, user threads would contend at system call

■ User Threads

- Managed and scheduled by the thread library
- May be either bound or unbound





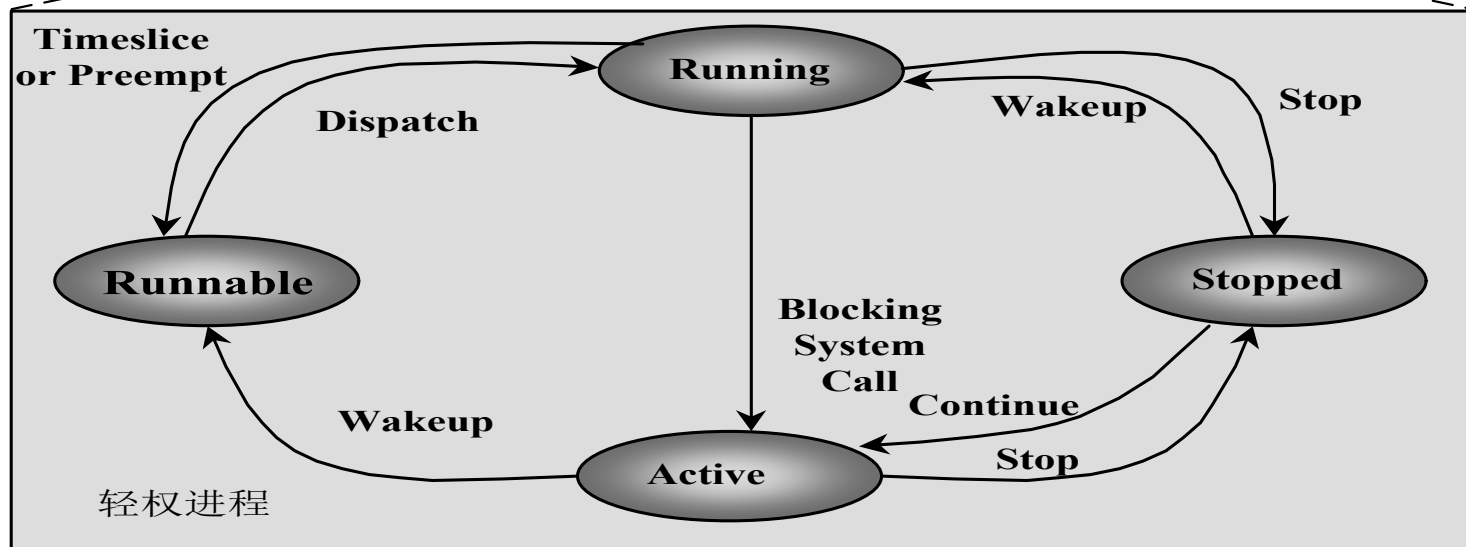
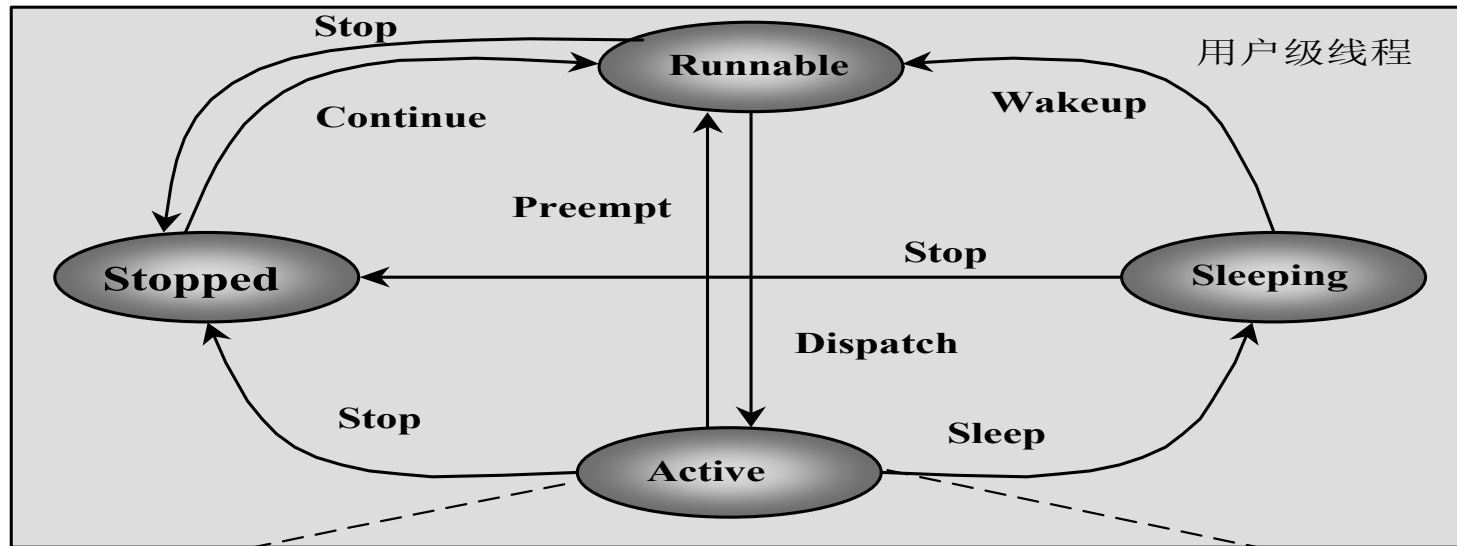
LWP

- LWP is supported by kernel, it is kernel data structure
- LWP is virtual execution environment for user threads.
 - When a user thread issues a system call, its register states are put in the stack of LWP
- Each process contains at least one LWP
- Each LWP is connected to exactly one kernel-level thread.





Solaris: User Threads and LWP





Resources required for different threads in Solaris

- Kernel thread: small data structure and a stack; thread switching relatively fast.
- LWP: a register set for the user-thread it is running, accounting and memory information, switching between LWPs is relatively slow.
- User-level thread: only need stack and program counter; no kernel involvement means fast switching.
- Kernel only sees the LWPs that support user-level threads.





Solaris threads

- User level threads may be either bound or unbound
- Bound
 - A user thread is permanently attached to a LWP
- All unbound threads in an application are multiplexed onto the pool of available for the application
 - Threads are unbound by default .
- The thread library adjusts LWPs in the pool
 - The thread library ages LWPs and deletes them when they are unused for a long time, typically about 5 minutes.



End of Chapter 4

