


# Chapter 7: Deadlocks

---





# Chapter 7: Deadlocks

---

- The Deadlock Problem
- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock





# Chapter Objectives

---

- To develop a description of deadlocks, which prevent sets of concurrent processes from completing their tasks
- To present a number of different methods for preventing or avoiding deadlocks in a computer system





# The Deadlock Problem





# The Deadlock Problem

---

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.
- The resources may be either physical resources or logical resources.





# The Deadlock Problem

## ■ Example

- System has 2 disk drives
- $P_1$  and  $P_2$  each hold one disk drive and each needs another one

## ■ Example

- semaphores  $A$  and  $B$ , initialized to 1

$P_0$

wait ( $A$ );

wait ( $B$ );

$P_1$


wait( $B$ )

wait( $A$ )





# Deadlock Example



(a) A potential deadlock

(b) an actual deadlock








# Why Deadlock?

---

- Resource competition
  - The number of resources is not sufficient for processes requirements
- Advance order of Concurrent processes is not appropriate
  - In a multiprogramming system, it is asynchronous for concurrent processes execution
  - In some order, each process can complete successfully
  - In some order, there will be a deadlock

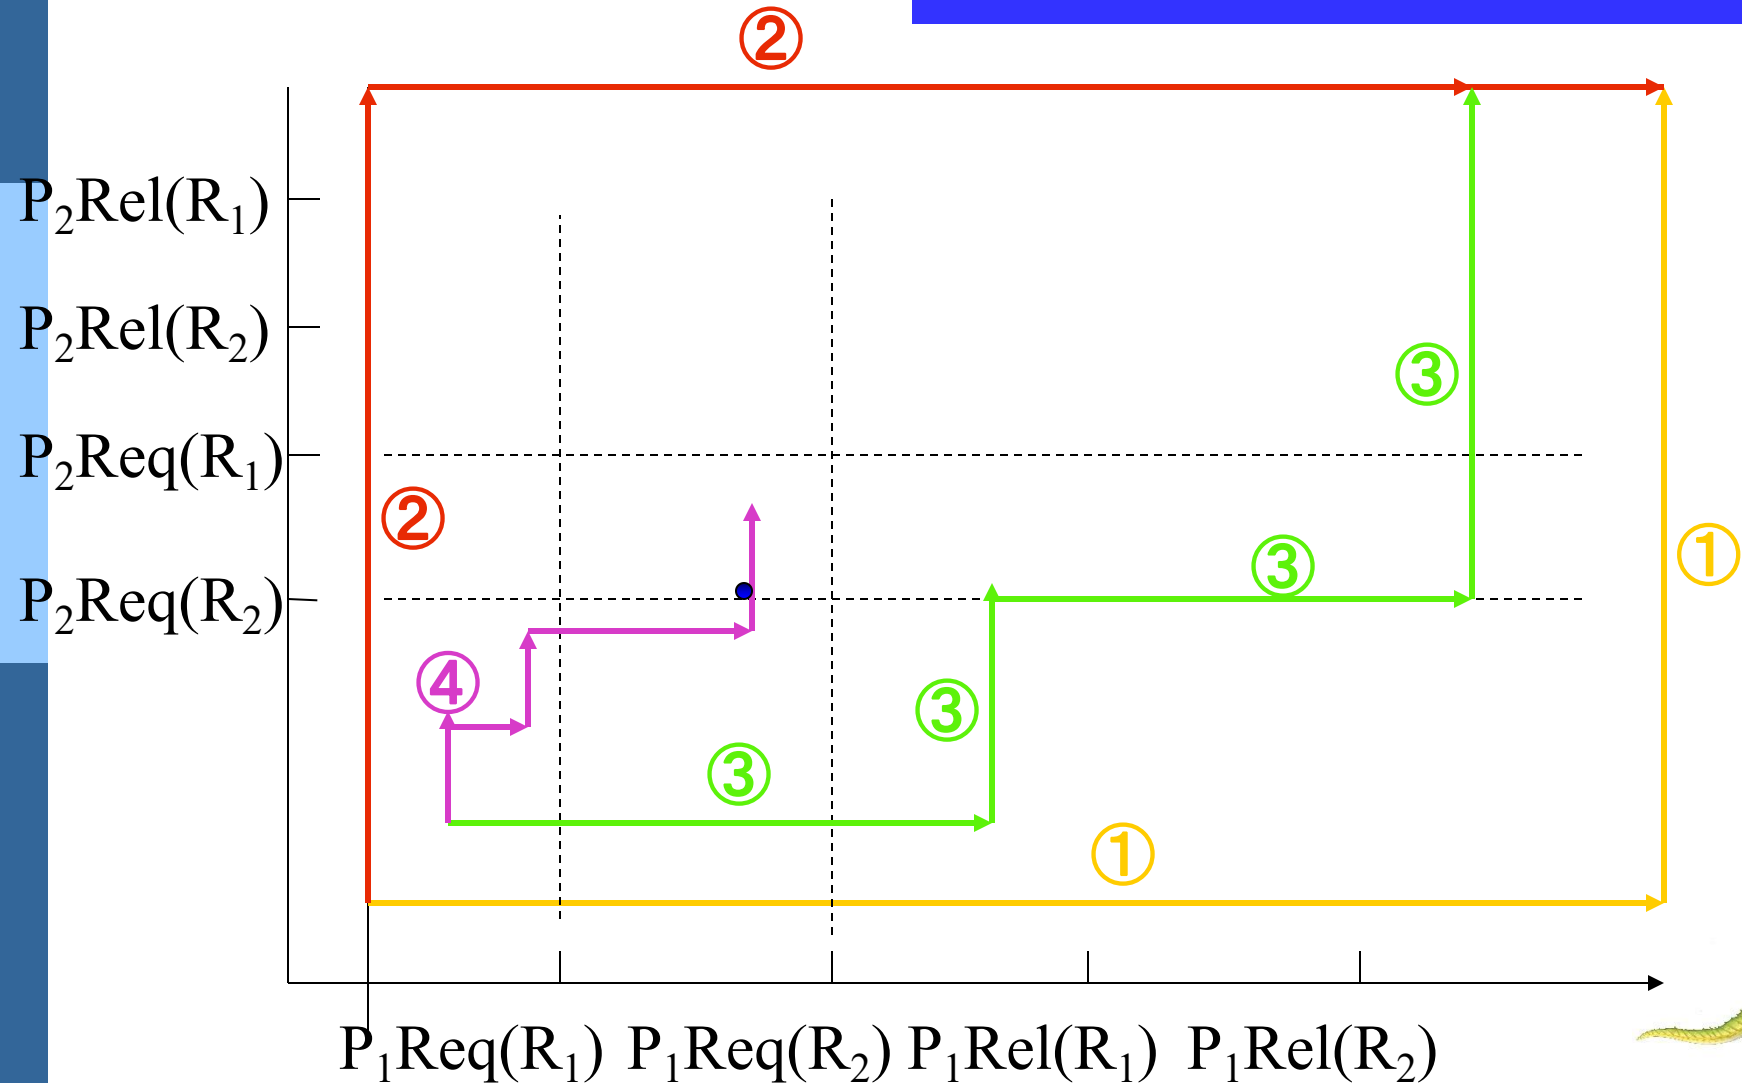






**Concurrent processes P1,P2; resources R1 and R2 has only one instance respectively**

**Curve4 will lead to unsafe area**





# System Model

---

- Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances.
- The number of resources requested may not exceed the total number of resources available in the system.





# How process utilizes a resource?

---

- Each process utilizes a resource as follows:
  - **request**
  - **use**
  - **Release**
- The request and release are system calls, such as `open()` and `close()` file, and `allocate()` and `free()` memory system calls.





# Necessary Conditions for Deadlock

Deadlock can arise if four conditions hold **simultaneously**.


- **Mutual exclusion:** only one process at a time can use a resource
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task





# Necessary Conditions for Deadlock(Cont)

- **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_{n-1}\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_0$ .





# Resource-Allocation Graph


A set of vertices  $V$  and a set of edges  $E$ .

- $V$  is partitioned into two types:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
- request edge – directed edge  $P_i \rightarrow R_j$
- assignment edge – directed edge  $R_j \rightarrow P_i$







# Resource-Allocation Graph



(a)



(b)




(c)

- ➔ Resource R assigned to process A
- ➔ Process B is requesting/waiting for resource S
- ➔ Process C and D are in deadlock over resources T and U







# Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances



- $P_i$  requests instance of  $R_j$




- $P_i$  is holding an instance of  $R_j$






# Example of a Resource Allocation Graph






# Resource Allocation Graph With A Deadlock





# Graph With A Cycle But No Deadlock





# Basic Facts

---

- If graph contains no cycles  $\Rightarrow$  no deadlock
- If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock
  - if several instances per resource type, possibility of deadlock





# Methods for Handling Deadlocks

---

- Ensure that the system will *never* enter a deadlock state
  - Deadlock prevention
  - Deadlock avoidance
- Allow the system to enter a deadlock state and then recover
  - Deadlock detection and recovery
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX
  - Ostrich policy





# Deadlock Prevention

- Restrain the ways request can be made, so that at least one of the necessary conditions for deadlock can not occur.
- Possible side effects:
  - Low device utilization
  - Reduced system throughput

## ■ Mutual Exclusion

- not required for sharable resources; must hold for nonsharable resources.







# Deadlock Prevention

---


- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
  - Require process to request and be allocated all its resources before it begins execution
  - or allow process to request resources only when the process has none
  - Low resource utilization; starvation possible





# Dining Philosophers

- 5 philosophers sitting around a round table
- 1 chopstick in between each pair of philosophers
  - 5 chopsticks total
- Each philosopher needs two chopsticks to eat
- How to prevent deadlock





# Deadlock Prevention (Cont.)

---

## ■ No Preemption –

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
- Preempted resources are added to the list of resources for which the process is waiting
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting





# Deadlock Prevention (Cont.)

---

## ■ No Preemption –

- When process A requests some unavailable resources, we check whether they are allocated to another process B that is waiting for other resources;
- If so, we preempt the desired resources from B and allocate them to A.
- Otherwise, A must wait.





# Deadlock Prevention (Cont.)

---


- Allow Preemption!
- Can preempt CPU by
  - saving its state to thread control block and resuming later
- Can preempt memory by
  - swapping memory out to disk and loading it back later
- Can we preempt the holding of a lock?





# Deadlock Prevention (Cont.)

- Some resource cannot be preempted
- Consider a process given the printer
  - halfway through its job
  - now forcibly take away printer
  - !!??





# Deadlock Prevention (Cont.)

---

## ■ Circular Wait

- impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration







# Deadlock Avoidance

---

- Requires that the system has some additional *a priori* information available. E.g, which type resources a process needs? How many instances of each resource type a process needs at most?
- Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need





# Deadlock Avoidance

---

- The deadlock-avoidance algorithm **dynamically** examines the resource-allocation state to ensure that there can never be a circular-wait condition
- **Resource-allocation state** is defined by the number of available and allocated resources, and the maximum demands of the processes
- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state.





# Safe State

- System is in **safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes is the systems such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$
- That is:
  - If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on





## example

- processes:  $P_1$ ,  $P_2$  and  $P_3$ . 12 tapes. Snapshot at  $T_0$  :

process	max	allocated	need	available
$P_1$	10	5	5	3
$P_2$	4	2	2	
$P_3$	9	2	7	

- sequence ( $P_2$ ,  $P_1$ ,  $P_3$ ) satisfies the safety condition :
  - Allocate 2 of the 3 available tapes to  $P_2$ , which satisfies the max need of  $P_2$ . when  $P_2$  finishes, the system will have 5 available tapes.
  - Then allocate the 5 tapes to  $P_1$ . when  $P_1$  finishes, the system will have 10 available tapes.
  - At last, allocate 7 of the 10 tapes to  $P_3$
  - **So, system is safe at  $T_0$ .**





## System goes from a safe state to an unsafe state

- Suppose that, at  $T_0$ ,  $P_3$  requests a tape and is allocated one tape. Then system goes from  $T_0$  to  $T_1$ .

- Process    max    allocated    need    available    available  
   before allocating    after releasing

$P_1$	10	5	5	>		
$P_2$	4	2	2	=<	2	4
$P_3$	9	3	6	>		

- At  $T_1$ , we can not find a sequence satisfying safety condition, so, it is unsafe.
- Therefore, when  $P_3$  requests a tape at  $T_0$ , we must make it wait to avoid deadlock.





# Basic Facts


---

- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.





# Safe, Unsafe , Deadlock State







# Avoidance algorithms

---

- Single instance of a resource type
  - Use a resource-allocation graph
  
- Multiple instances of a resource type
  - Use the banker's algorithm






# Resource-Allocation Graph Scheme

- **Claim edge**  $P_i \rightarrow R_j$  indicated that process  $P_j$  may request resource  $R_j$ ; represented by a dashed line
- Claim edge converts to **request edge** when a process requests a resource
- Request edge converted to an **assignment edge** when the resource is allocated to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system






# Resource-Allocation Graph





# Unsafe State In Resource-Allocation Graph





# Resource-Allocation Graph Algorithm

---

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph





# Banker's Algorithm

---

- Multiple instances
- Each process must claim the maximum number of each resource type it may need
- When a process requests a resource it may have to wait
- When a process gets all its resources it must return them in a finite amount of time





# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes,

$m$  = number of resources types.

- **Available:** Vector of length  $m$ . If available  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max:**  $n \times m$  matrix. If  $Max[i, j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- **Allocation:**  $n \times m$  matrix. If  $Allocation[i, j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need:**  $n \times m$  matrix. If  $Need[i, j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task

$$Need[i, j] = Max[i, j] - Allocation[i, j]$$





# Safety Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:

*Work* = *Available*

*Finish* [*i*] = *false* for *i* = 0, 1, ..., *n* - 1

2. Find an *i* such that both:

(a) *Finish* [*i*] = *false*

(b)  $Need_i \leq Work$

If no such *i* exists, go to step 4

3. *Work* = *Work* + *Allocation*<sub>*i*</sub>

*Finish*[*i*] = *true*

go to step 2

4. If *Finish* [*i*] == *true* for all *i*, then the system is in a safe state







# Resource-Request Algorithm for Process $P_i$

$Request_i$  = request vector for process  $P_i$ . If  $Request_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If  $Request_i \leq Need_i$  go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If  $Request_i \leq Available$ , go to step 3. Otherwise  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If safe  $\Rightarrow$  the resources are allocated to  $P_i$
- If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored





# Example of Banker's Algorithm

■ 5 processes  $P_0$  through  $P_4$ ;

3 resource types:

A (10 instances), B (5 instances), and C (7 instances)

Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	





## Example (Cont.)

- The content of the matrix *Need* is defined to be *Max – Allocation*

	<u>Need</u>		
	A	B	C
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1





## Example (Cont.)

- To find a safe sequence using safety algorithm

- $Work[] = available = (3, 3, 2)$

- $Finish[i] = false \ (i = 0..4)$

- 

- Work    need    allocation    finish

- P1   3 3 2   1 2 2   2 0 0      T


- P3   5 3 2   0 1 1   2 1 1      T

- P4   7 4 3   4 3 1   0 0 2      T

- P2   7 4 5   6 0 0   3 0 2      T

- P0   10 4 7   7 4 3   0 1 0      T

- The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria





## Example (Cont.)

- Usually, there would be more than one safe sequences if the system is in safe state.

- $Work[] = available = (3, 3, 2)$

- $Finish[i] = false \quad (i = 0..4)$

- 

- 

	Work	need	allocation	finish
--	------	------	------------	--------

- P3    3 3 2        0 1 1        2 1 1        T

- P4    5 4 3        4 3 1        0 0 2        T

- P1    5 4 5        1 2 2        2 0 0        T

- P2    7 4 5        6 0 0        3 0 2        T

- P0    10 4 7        7 4 3        0 1 0        T

- Safe sequence: (P3, P4, P1, P2, P0)





## Example: $P_1$ Request (1,0,2)

- Check that Request  $\leq$  Available (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true)

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	4	3	2	3	0
$P_1$	3	0	2	0	2	0			
$P_2$	3	0	1	6	0	0			
$P_3$	2	1	1	0	1	1			
$P_4$	0	0	2	4	3	1			

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- Can request for (3,3,0) by  $P_4$  be granted?
- Can request for (0,2,0) by  $P_0$  be granted?





# Deadlock Detection

---

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme





# Single Instance of Each Resource Type


- Maintain *wait-for* graph
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph








# Resource-Allocation Graph and Wait-for Graph



(a)

Resource-Allocation Graph



(b)

Corresponding wait-for graph





# Several Instances of a Resource Type

---

- **Available:** A vector of length  $m$  indicates the number of available resources of each type.
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If  $Request[i, j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .





# Detection Algorithm

---

1. Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively Initialize:

(a) *Work* = *Available*

(b) For  $i = 1, 2, \dots, n$ , if  $Allocation_i \neq 0$ , then  
 $Finish[i] = false$ ; otherwise,  $Finish[i] = true$

2. Find an index  $i$  such that both:

(a)  $Finish[i] == false$

(b)  $Request_i \leq Work$

If no such  $i$  exists, go to step 4





## Detection Algorithm (Cont.)

3.  $Work = Work + Allocation_i$   
 $Finish[i] = true$   
go to step 2
4. If  $Finish[i] == false$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if  $Finish[i] == false$ , then  $P_i$  is deadlocked

**Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state**






# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>			<u>Request</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	0			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

- Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$





## Example (Cont.)

■  $P_2$  requests an additional instance of type C

Allocation    Request    Available

	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	1			
$P_2$	3	0	3	0	0	1			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

■ State of system?

- Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes requests
- Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$





# Detection-Algorithm Usage

---

- When, and how often, to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - ▶ one for each disjoint cycle





# Detection-Algorithm Usage

---

- Deadlock occurs only when some process makes a request that can not be granted immediately.
- We can invoke deadlock detection algorithm every time a resource request cannot be granted immediately.
  - Lead to a considerable overhead in computation time.
- A less expensive alternative is to invoke the algorithm at less frequent intervals.
  - Once per hour
  - Whenever CPU utilization drops below 40%.







# Recovery from Deadlock: Process Termination

---

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated
- In which order should we choose to abort?
  - Priority of the process
  - How long process has computed, and how much longer to completion
  - Resources the process has used
  - Resources process needs to complete
  - How many processes will need to be terminated
  - Is process interactive or batch?





# Recovery from Deadlock: Resource Preemption

---

- Selecting a victim – minimize cost
- Rollback – return to some safe state, restart process for that state
- Starvation – same process may always be picked as victim, include number of rollback in cost factor





# assignment

---

- 7.1
- 7.9
- 7.13



# End of Chapter 7

---

