


# Chapter 11: File System Implementation

---





# Chapter 11: File System Implementation

---

- File-System Structure
- File-System Implementation
- Directory Implementation
- Allocation Methods
- Free-Space Management
- Efficiency and Performance
- Recovery





# Objectives

---

- To describe the details of implementing local file systems and directory structures
- To describe the implementation of remote file systems
- To discuss block allocation and free-block algorithms and trade-offs





# File-System Structure

---

- A file system includes:
  - File :Logical storage unit
    - ▶ Collection of related information
  - Directory
    - ▶ **File control block** – storage structure consisting of information about a file
  - Software





# File-System Structure

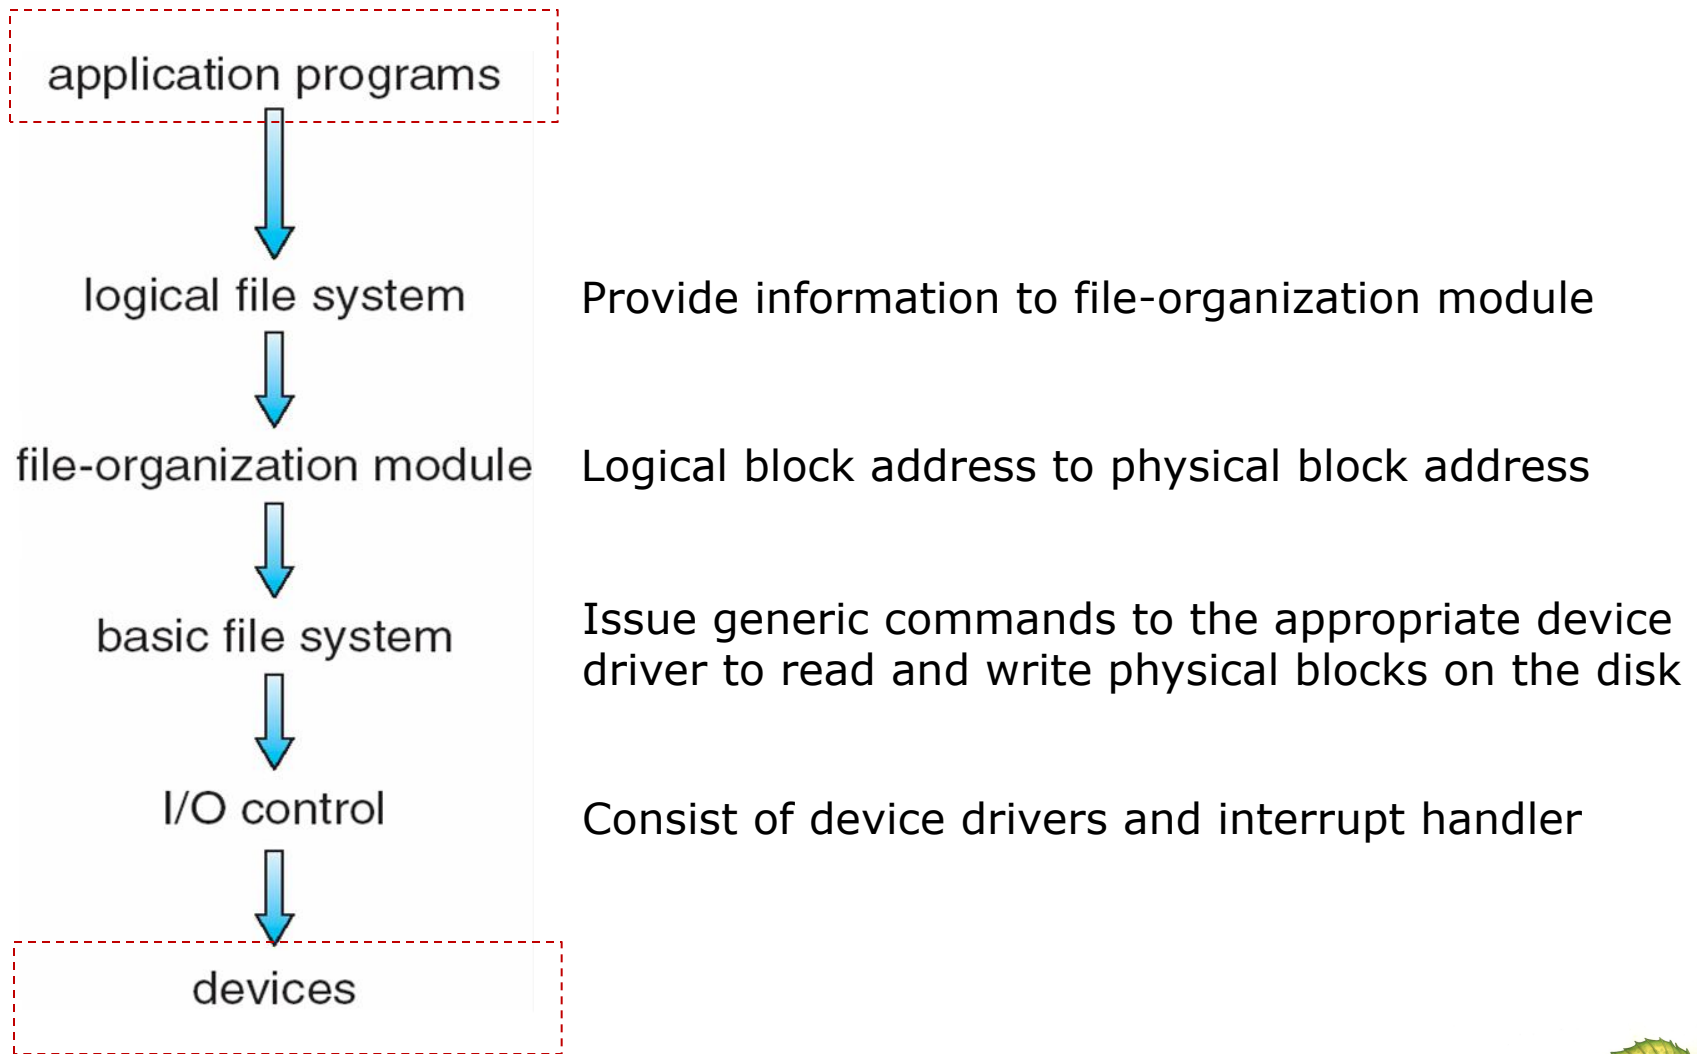
---

- **File system** resides on secondary storage (disks)
  - Provides efficient and convenient access to disk by allowing data to be stored, located , retrieved easily
- File system organized into layers.
- **Device driver** controls the physical device
  - Transfer information between the main memory and the disk.





# Layered File System



管理文件目录,根据文件名得到该文件的相关信息,提供给下一层;文件的保护和安全

找到I/O设备;实现逻辑记录到数据块的映射;磁盘调度,性能优化

## 件集合

## 基本文件系统（物理 I/O 层）

## 向驱动程序发出读/写数据块的命令

# device



# File-System Implementation

- Several on-disk and in-memory structures are used to implement a file system
- On-disk structures includes:
  - **Boot control block** contains info needed by system to boot OS from that volume
  - **Volume control block** contains volume details
    - ▶ 分区块的数量、块大小、空闲块数量和指针、空闲的FCB数量和指针
    - ▶ UNIX FS: Super block (超级块)
    - ▶ NTFS: Master file table (主控文件表)
  - Directory structure organizes the files
  - Per-file **File Control Block (FCB)** contains many details about the file







# A Typical File Control Block

---

file permissions

file dates (create, access, write)

file owner, group, ACL

file size

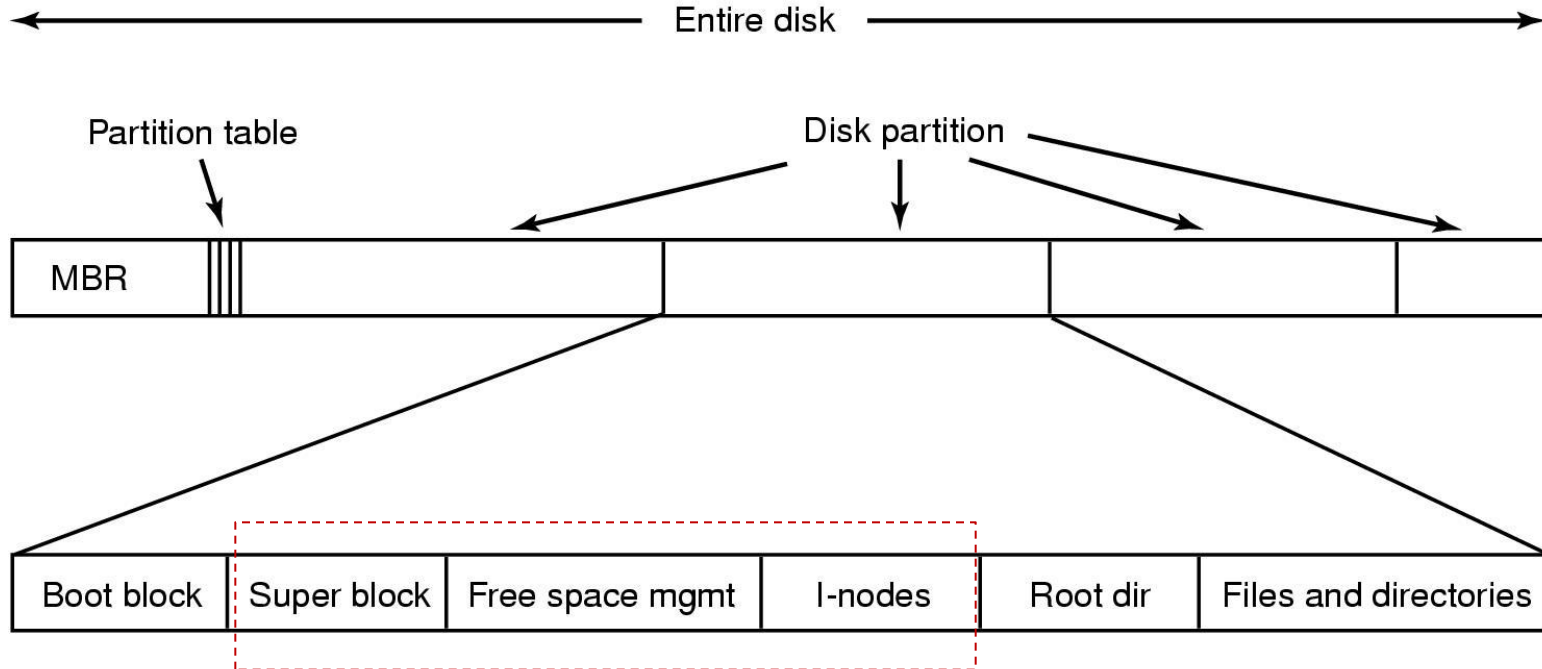
file data blocks or pointers to file data blocks





# File System Layout

## Unix file system layout



## Window FAT file system layout





# File System Layout

---

- Sector 0 of a disk is called MBR
  - Master Boot Record (主引导记录)
- MBR is used to boot the computer
- The end of MBR contains the partition table
  - Giving starting and ending addresses of each partition
  - One of the partitions is marked as active in the table





# File System Layout

---


- At boot time, BIOS reads in the MBR
- MBR then locates the active partition
  - and reads in the first block, the boot block (引导块), and execute it
- Boot block in turns loads the OS in the partition





# File System Layout

- boot block后面的内容就因系统而异
- boot block后面是一个超级块(分区控制块)
  - 存放该文件系统的各种参数:文件系统类型, 数据块尺寸, 空闲块的数量和指针, 空闲FCB的数量和指针等等
- 超级块后面是磁盘自由空间
- 再后面是I-NODE区 (UNIX), FAT文件系统无此区
- 再后面是根目录区
- 最后是其他目录和用户文件区





# In-Memory File System Structures

---


- The in-memory structures may include the ones as below:
  - An **mount-table**: contains info about each mounted volume
  - An **directory-structure cache**: holds the directory info of recent accessed directories
  - The **system-wide open-file table**
  - The **per-process open-file table**





## ■ System-wide open-file table

- 整个系统一张
- 放在内存：用于保存已打开文件的FCB



## ■ Per-process open-file table

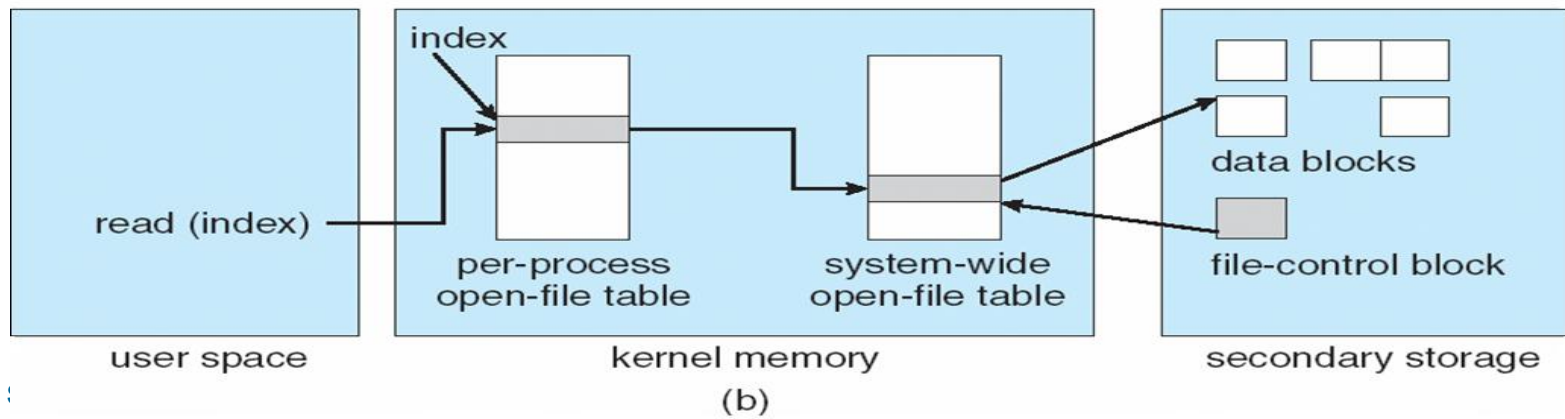

- 每个进程一张
- 进程的FCB中记录了用户打开文件表的位置





# In-Memory File System Structures

- The following figure illustrates the necessary file system structures provided by the operating systems.
  - Figure (a) refers to opening a file.
  - Figure (b) refers to reading a file.







# Virtual File Systems

---

- Modern OS must support multiple type of file systems.
  - How to allow multiple types of file system to be integrated into a directory structure?
  - How to support users' seamless move between different file systems?
- Most OS, including UNIX, use object-oriented techniques to simplify, organize and modularize the implementation.






# Virtual File Systems

---


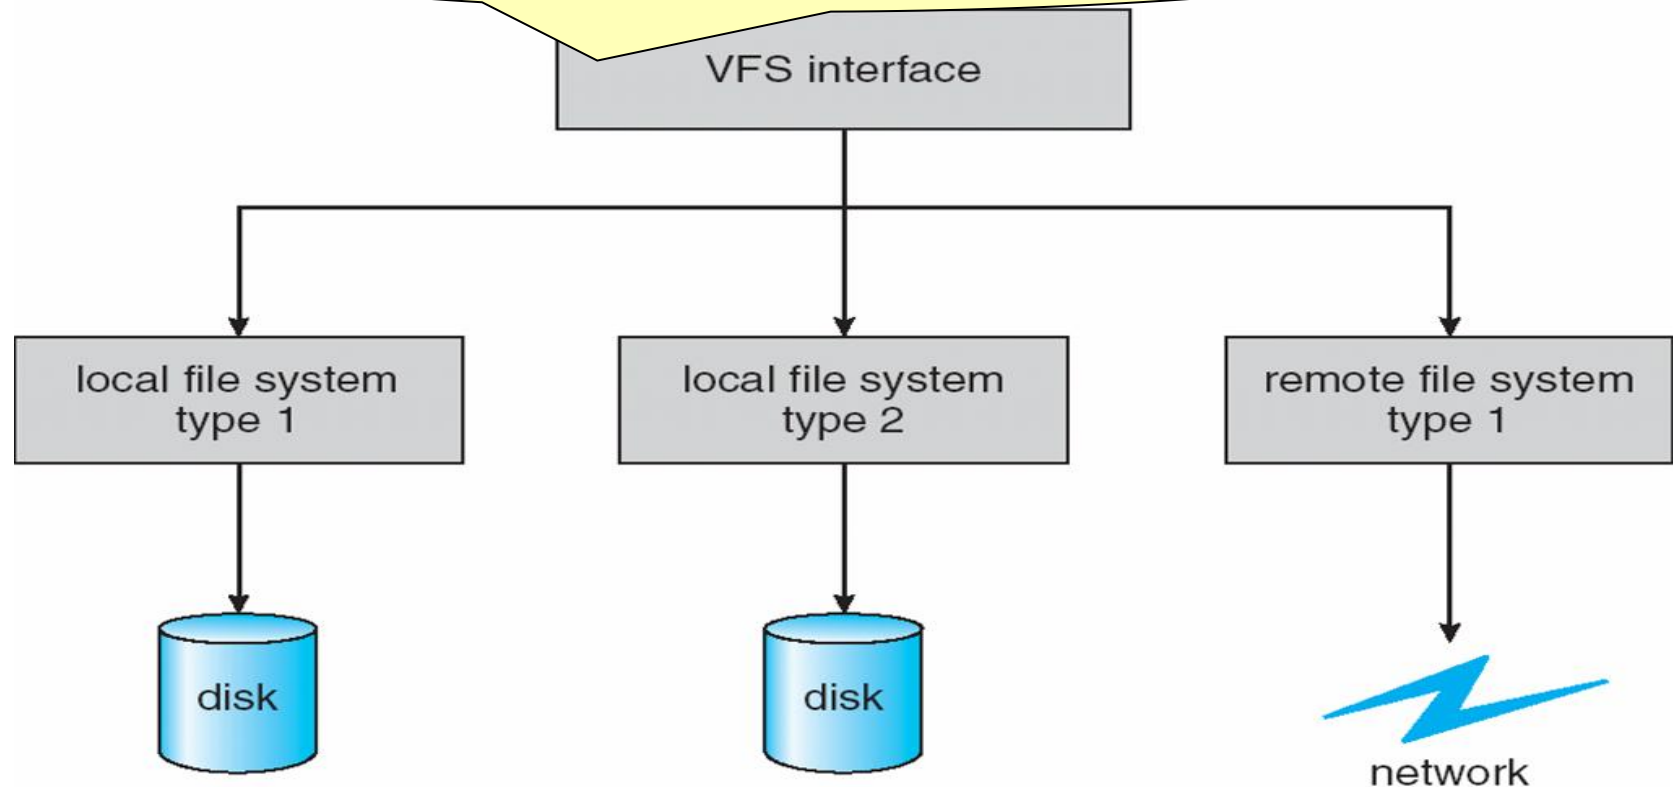
- Virtual File Systems (VFS) provide an object-oriented way of implementing file systems.
- VFS allows the same system call interface (the API) to be used for different types of file systems.
- The API is to the VFS interface, rather than any specific type of file system.





Including open(), read(), write(), and close() calls and file descriptors.

Separating standard file-system Operations from their implementation ; providing a mechanism for uniquely representing a file throughout a network based on **vnode**





# Directory Implementation

- The selection of directory-allocation and directory-management algorithms significantly affects the efficiency, performance and reliability of the file system.
- **Linear list** (线性列表) of file names with pointer to the data blocks.
  - simple to program
  - time-consuming to search
- **Hash Table** – linear list with hash data structure.
  - decreases directory search time
  - **collisions** – situations where two file names hash to the same location
  - fixed size

FCB

**Name**

**Location**

**Size**

...





# 文件目录改进

如何加快目录检索？

- 为加快目录检索可采用目录项分解法：把FCB分成两部分：
  - 符号目录项：
    - 文件名，文件号(iNode号)
  - 基本目录项（索引节点目录iNode）：
    - 除文件名外的所有字段






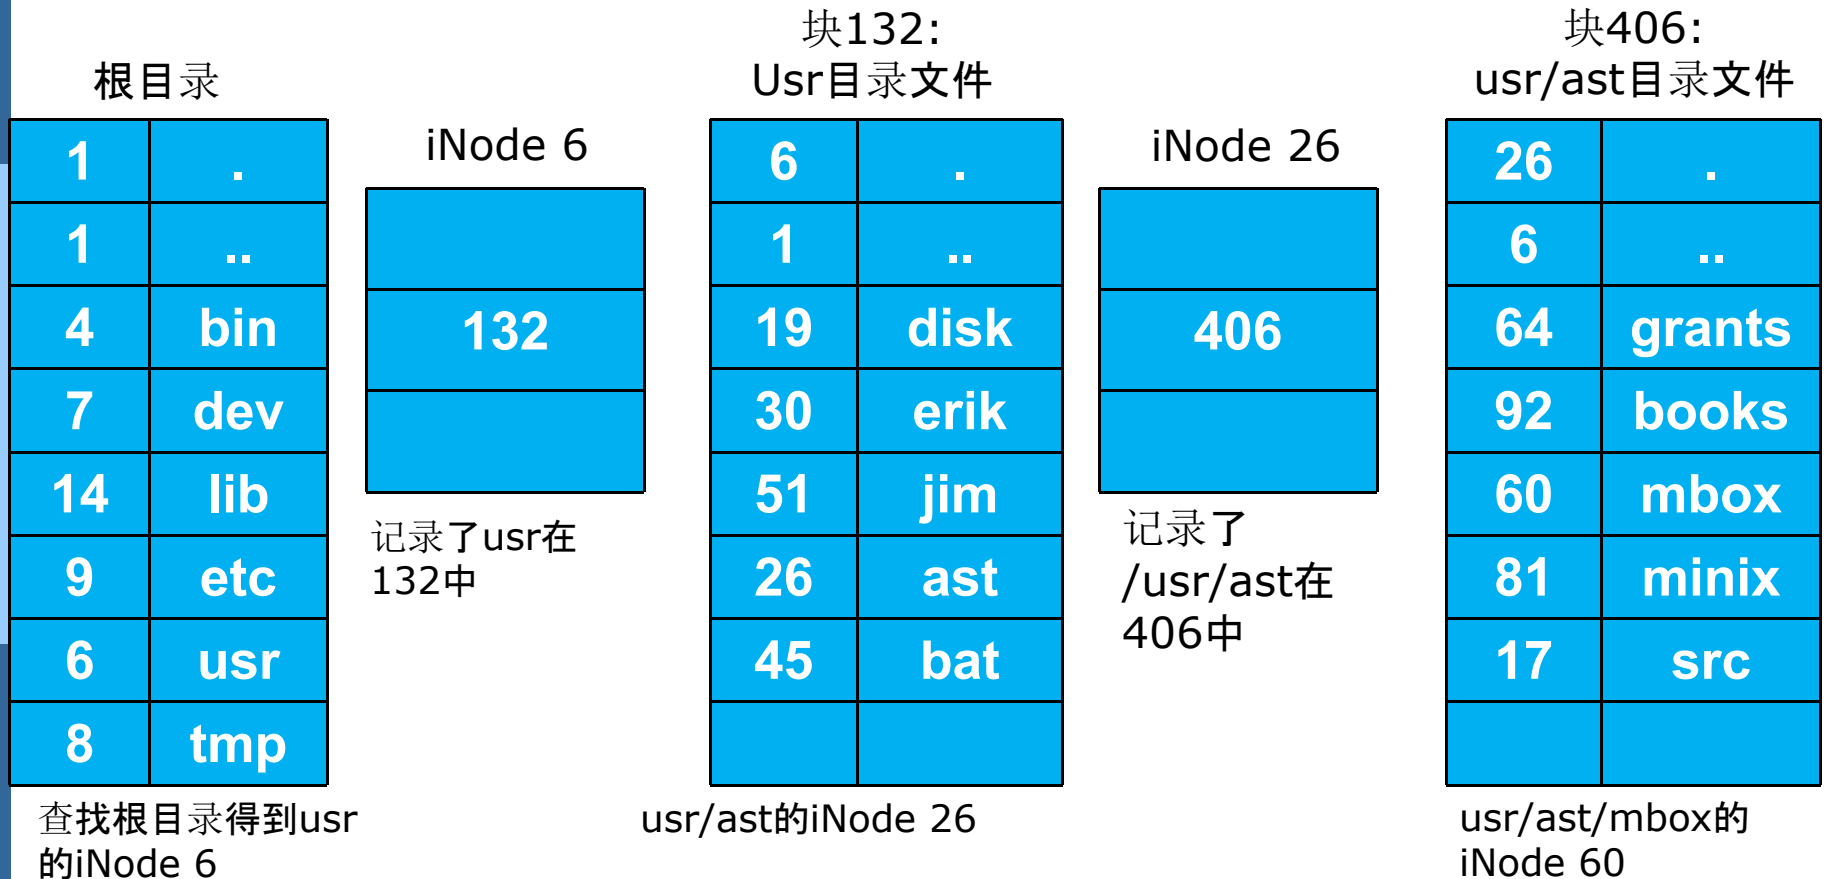
# UNIX的目录结构

- UNIX采用文件名和文件说明分离的目录结构
  - 每个文件有一个存放在磁盘索引节点区的索引节点，称为**磁盘索引节点**，它包括以下内容：
    - 文件主标识符和同组用户标识符；
    - 文件类型：是普通文件、目录文件、符号连接文件或特别文件（又分块设备文件或字符设备文件）
    - 文件主人，同组用户和其它人对文件存取权限（读R、写W、执行X）
    - 文件的物理地址，用于UNIX直接、间接混合寻址的13个地址项 `di_addr[13]`
    - 文件长度（字节数） `di_size`
    - **文件链接数** `di_nlink`
    - 文件最近存取和修改时间等。





## ■ 给定文件路径名为 /usr/ast/mbox，检索过程如下





# 习题

■ 假设 一个FCB 占64个字节，物理块大小512字节。

符号目录项占10字节（文件名8字节，文件号2字节）。

基本目录项（iNode）占 $64-8=56$ 字节

若一个目录文件有254个目录项，请分别给出分解前后查找文件的某一个FCB 的平均磁盘访问次数。

解答：

分解前：占 $254*64/512=32$ 块

分解后：符号文件占 $254*10/512=5$ 块

查找一个文件的FCB的平均访盘次数：

分解前：16.5次 =  $(1+2+3+\dots+32) / 32$

分解后：4次 =  $(2+3+4+5+6) / 5$







# 文件的物理结构

- 文件的物理结构
  - 文件在存储介质上的存放方式
  - 主要解决两个问题：
    - 假设一个文件被划分为N块，这N块在磁盘上是怎么存放的？
    - 其地址（块号）在FCB中是怎么记录的？
- 连续结构
- 链接结构
- 索引结构





# Allocation Methods

---

- An allocation method refers to how disk blocks are allocated for files:
- Contiguous allocation
- Linked allocation
- Indexed allocation





# Contiguous Allocation

---


- Each file occupies a set of contiguous blocks on the disk
- Simple – only starting location (block #) and length (number of blocks) are required
- Support sequential access and direct access
- Wasteful of space (dynamic storage-allocation problem) 碎片？ 外碎片
- Files cannot grow





# Contiguous Allocation

- Mapping from logical to physical

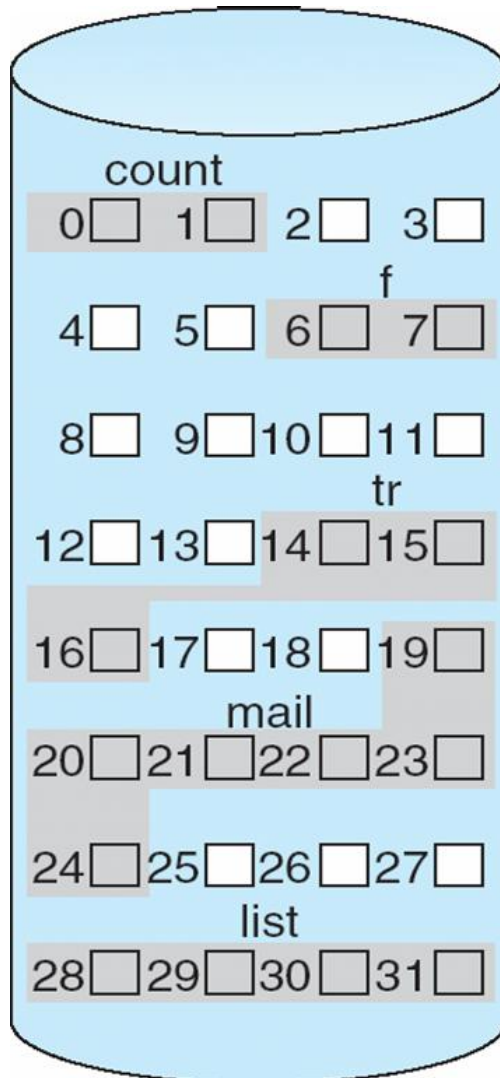


Block to be accessed =  $Q + \text{starting address}$   
Displacement into block =  $R$





# Contiguous Allocation of Disk Space



directory

| file  | start | length |
|-------|-------|--------|
| count | 0     | 2      |
| tr    | 14    | 3      |
| mail  | 19    | 6      |
| list  | 28    | 4      |
| f     | 6     | 2      |

FCB中如何记录文件地址?

第一块的块号 + 长度





# Contiguous Allocation

## ■ 优点

- 简单
- 支持顺序存储和随机存取
- 所需的磁盘寻道次数和寻道时间最少
- 可以同时读入多个块，检索一个块也很容易

## ■ 缺点

- 文件不能动态增长
  - ▶ 预留空间：浪费      或      重新分配和移动
- 不利于文件插入和删除
- 外部碎片：紧缩技术





# Extent-Based Systems


- Many newer file systems (I.e. Veritas File System) use a modified contiguous allocation scheme
- Extent-based file systems allocate disk blocks in extents
  - Initially, a contiguous chunk of space is allocated to the file
  - Then, if the amount proves to be not enough, another chunk of contiguous space—an extent—is added
  - The location of the file is recorded as **a location and a block count, plus a link to the first block of the extent**
- An **extent** is a contiguous block of disks
  - Extents are allocated for file allocation
  - A file consists of one or more extents





# Linked Allocation


- Each file is a linked list of disk blocks: blocks may be scattered anywhere on the disk.
- The directory contains a pointer to the first and last blocks of the file.







# Linked Allocation



FCB中如何记录文件地址?


记录起始块号





# Linked Allocation (Cont.)

- Simple – need only starting address
- Free-space management system – no waste of space
- No random access
- Mapping



Block to be accessed is the Qth block in the linked chain of blocks representing the file.

Displacement into block =  $R + 1$





# Linked Allocation

## ■ 优点

- 提高了磁盘空间利用率，不存在外碎片问题
- 有利于文件插入和删除
- 有利于文件动态扩充

## ■ 缺点

- 存取速度慢，不适于随机存取
- 可靠性问题，如指针出错
- 更多的寻道次数和寻道时间
- 链接指针占用一定的空间






# File-Allocation Table

链接结构的一个变形

File-allocation table (FAT) :

disk-space allocation used by MS-DOS and OS/2.



思想:

将指针集中存放到一张表里  
---文件分配表 (FAT)

每个磁盘块都有一个条目;  
起始块号记录在FCB中。

最坏情况, 每块需要移动两次磁头。

FAT表项的值有3种:  
0, 下一块块号, -1

0表示空闲物理块

|     |        |        |     |         |
|-----|--------|--------|-----|---------|
| 引导区 | 文件分配表1 | 文件分配表2 | 根目录 | 其他目录和文件 |
|-----|--------|--------|-----|---------|





# Indexed Allocation


- Brings all pointers together into the **index block**
- Each file has its own index block, which is an array of disk-block addresses.

索引表就是磁盘块地址数组，其中第*i*个条目指向文件的第*i*块

- **The directory contains the address of the index block**
- **Logical view**

FCB中如何记录文件地址？

索引表存放在何处？

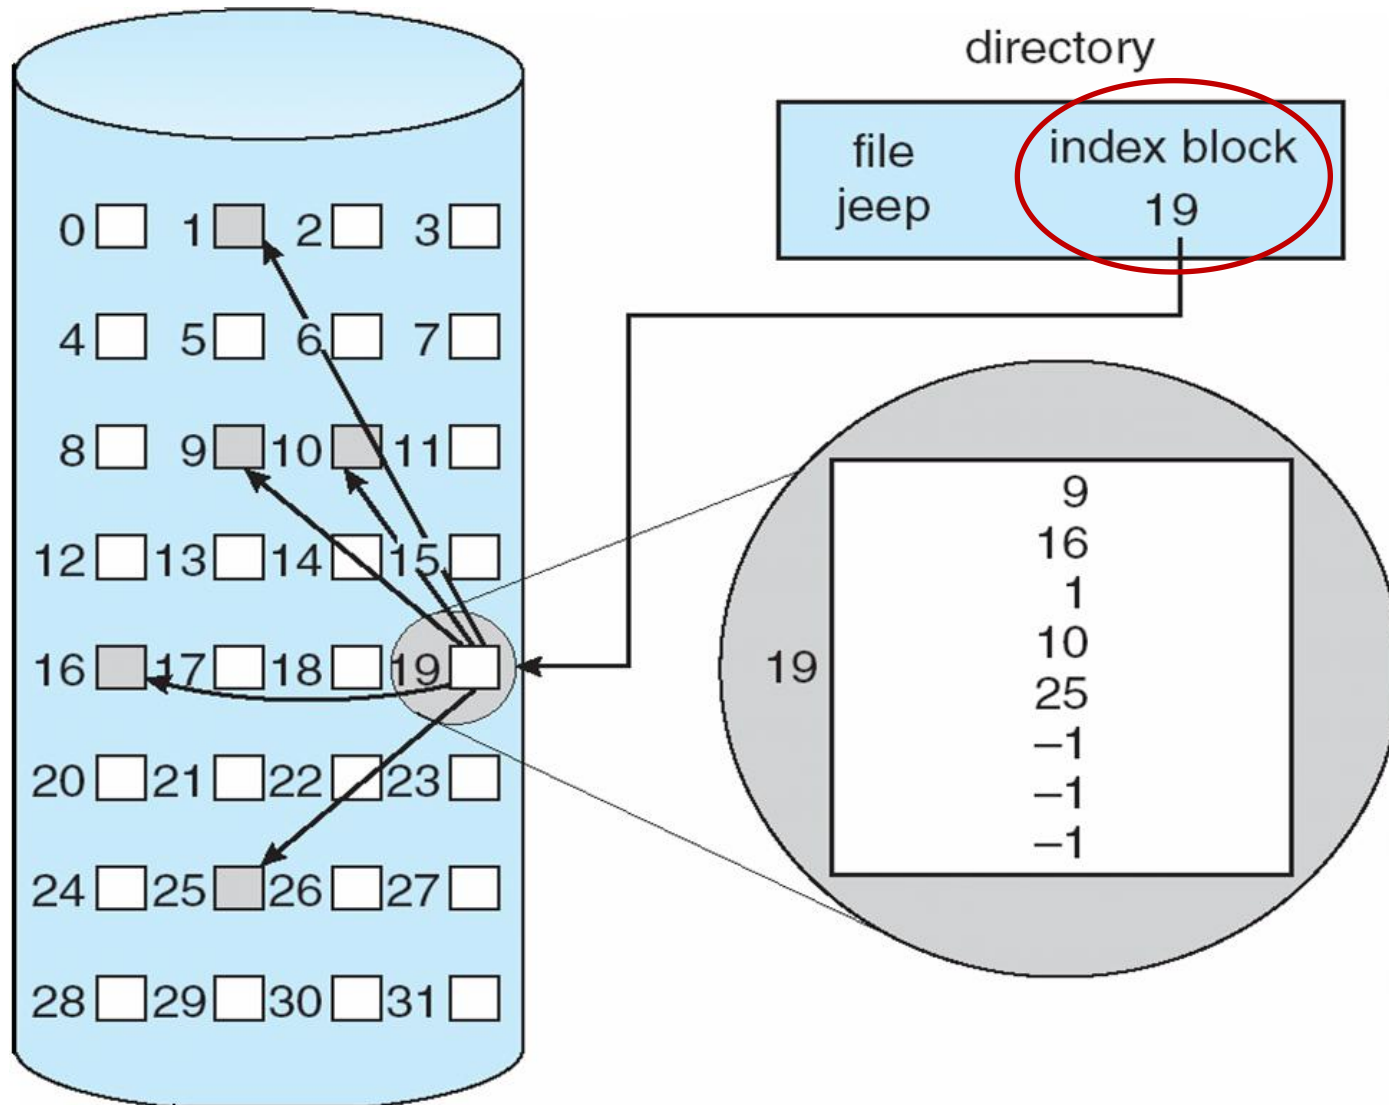


index table






# Example of Indexed Allocation





# Indexed Allocation (Cont.)

- Need index table
- Random access
- No external fragmentation, but have overhead of index block
- Mapping



Q = displacement into index table

R = displacement into block





# Indexed Allocation

## ■ 优点

保持链接结构的优点，又克服了其缺点：

- 既能顺序存取，又能随机存取
- 满足了文件动态增长、插入删除要求
- 能充分利用外存空间

## ■ 缺点

- 较多的寻道次数和寻道时间
- 索引表本身带来了系统开销, 如： 内外存空间，存取时间








# Indexed Allocation (Cont.)


---

- A large file may have **several index blocks**, how to organize them? (block size of 512 words)
  - Linked scheme – Link blocks of index table (no limit on size)
  - Two-level index (maximum file size is  $512^3$ )





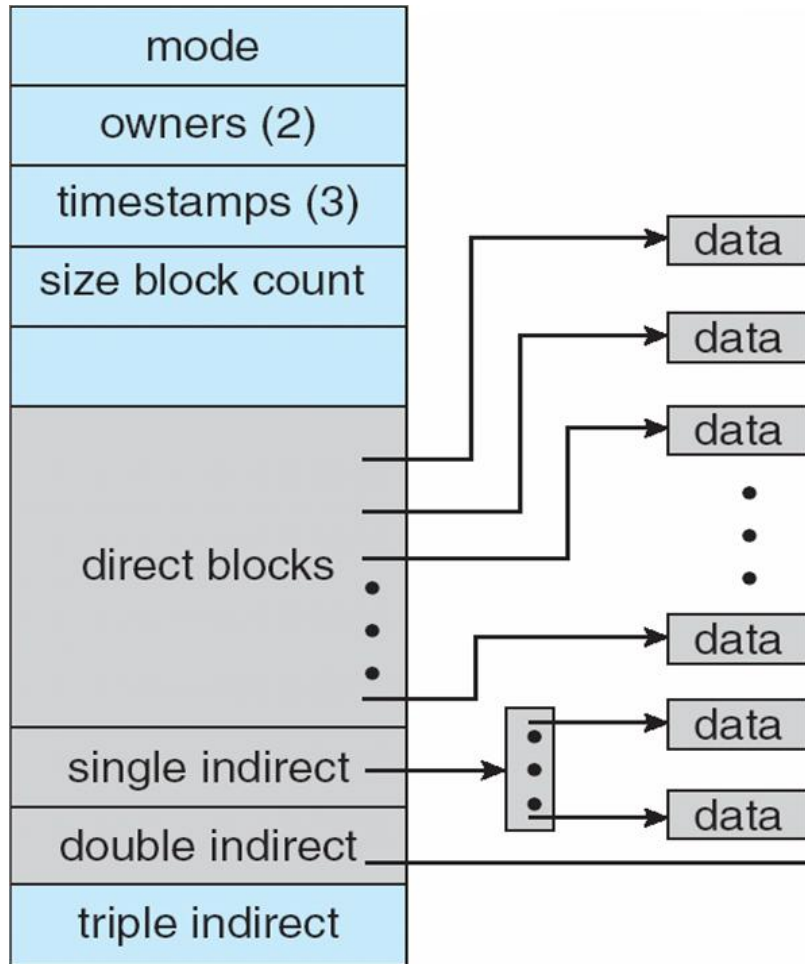
# Indexed Allocation – Mapping (Cont.)





# Combined Scheme: UNIX UFS (4K bytes per block)

## UNIX三级索引结构




iNode

- 每个文件的主索引表有**15**个索引项（FCB中），每项**2**个字节
- 前**12**项直接存放文件的物理块号
- 若文件大于**12**块，启用一级索引表（第**13**项）
- 第**14**项（二级索引表）
- 第**15**项（三级索引表）

假设扇区大小为512字节，物理块等于扇区大小

- 一级索引表可以存放?个物理块号 **256**
- 一个文件最大可达到?个物理块


$$12 + 256 + 256^2 + 256^3$$





# Free-Space Management

## ■ Bit vector ( $n$ blocks)



$\text{bit}[i] = \begin{cases} 1 \Rightarrow \text{block}[i] \text{ free} \\ 0 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$

## ■ Advantage

- Simple
- Efficient in finding the first free block or  $n$  consecutive free blocks

## ■ How to find the first free block

- The first non-0 word is scanned for the first 1-bit, which is the location of the first free block.

***(number of bits per word) \* (number of 0-value words) + offset of first 1-bit***

第一个空闲块号码（按字扫描）





# Free-Space Management (Cont.)

---

- Bit map requires extra space

- Example:

block size =  $2^{12}$  bytes

disk size =  $2^{30}$  bytes (1 gigabyte)

$n = 2^{30}/2^{12} = 2^{18}$  bits (or 32K bytes)






# Free-Space Management (Cont.)

- Linked list (free list) 空闲块链表
  - Cannot get contiguous space easily
  - No waste of space
- Grouping 成组链接法
  - 在第一个空闲块中存储 $n$ 个空闲块的地址
  - 前 $n-1$ 为空，第 $n$ 个指向另外 $n$ 个空闲块的地址
- Counting 空闲块表
  - 空闲空间表中每个条目记录一组连续空闲块的起始地址和数量.






# Linked Free Space List on Disk





# UNIX Grouping







# Efficiency and Performance

- 磁盘服务 --> 速度是系统性能的主要瓶颈之一
- 设计文件系统应尽可能减少磁盘访问次数
- 提高文件系统性能的方法：
  - 目录项（FCB）分解、当前目录、磁盘碎片整理、合理分配磁盘空间
  - 块高速缓存、提前读取
  - 磁盘调度、RAID技术等





# Efficiency and Performance

---

## ■ Efficiency dependent on:

- disk allocation and directory algorithms
- types of data kept in file's directory entry


## ■ Performance

- **disk cache** (文件缓存、磁盘高速缓存、块高速缓存、缓冲区高速缓存) – separate section of main memory for frequently used blocks 位于内核内存中
- **read-ahead** (提前读取) and **free-behind** (随后释放) – techniques to **optimize sequential access**





# Various Disk-Caching Locations





# Page Cache

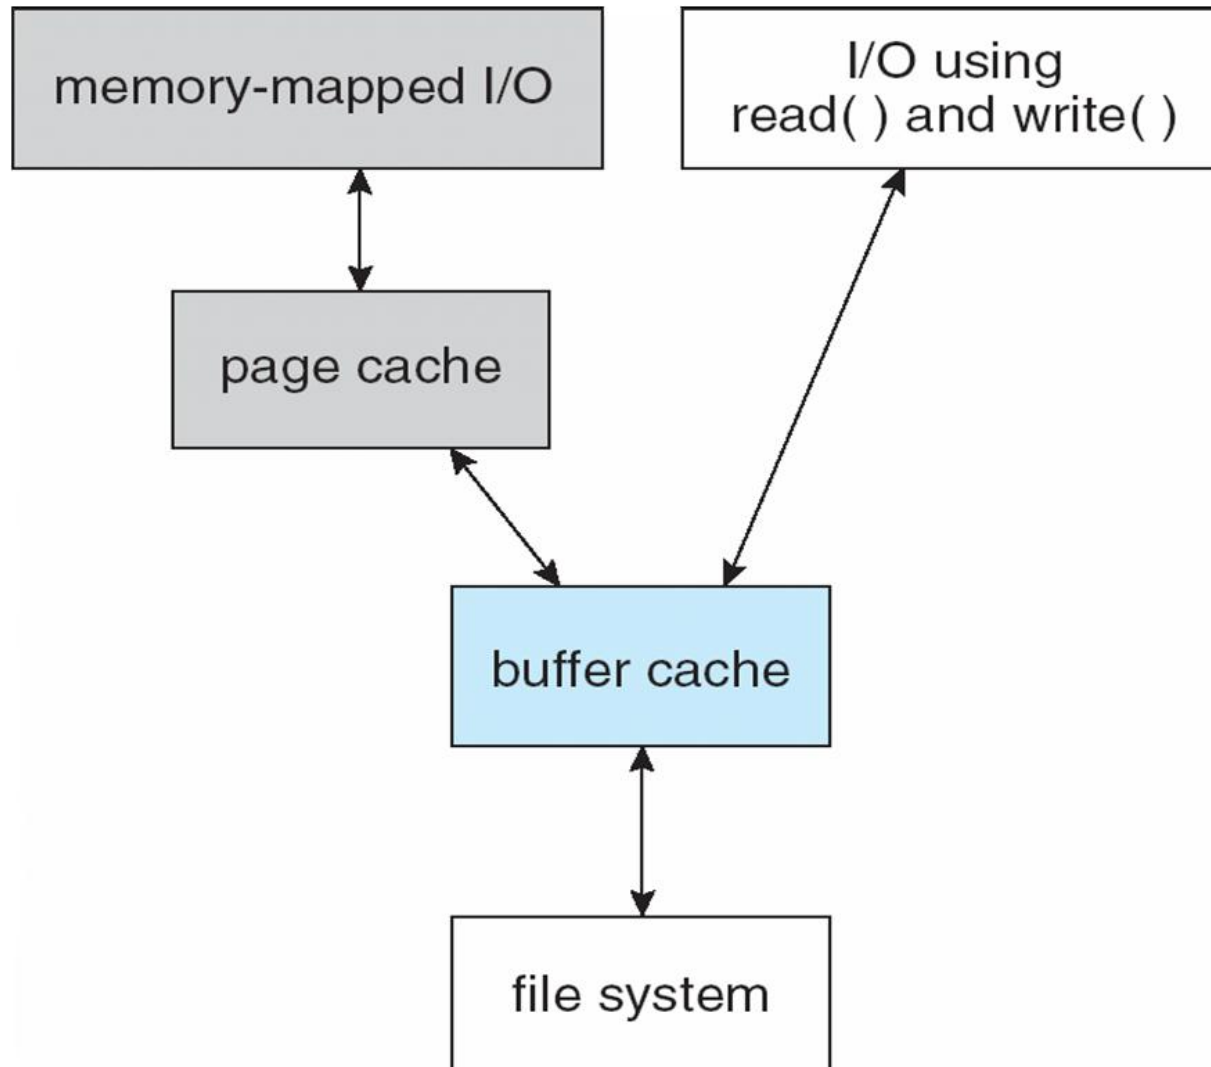
---

- A **page cache** caches pages rather than disk blocks using virtual memory techniques
- Memory-mapped I/O uses a page cache
- Routine I/O through the file system uses the buffer (disk) cache
- This leads to the following figure





# I/O Without a Unified Buffer Cache





# I/O Without a Unified Buffer Cache

---

- Memory-mapped I/O example:
  - Read in disk blocks from the file system and store them in the buffer cache;
  - Copy the blocks into the page cache
- This is called **double caching**
  - Waste memory
  - Waste CPU and I/O cycles
  - Inconsistencies between the two caches





# Unified Buffer Cache

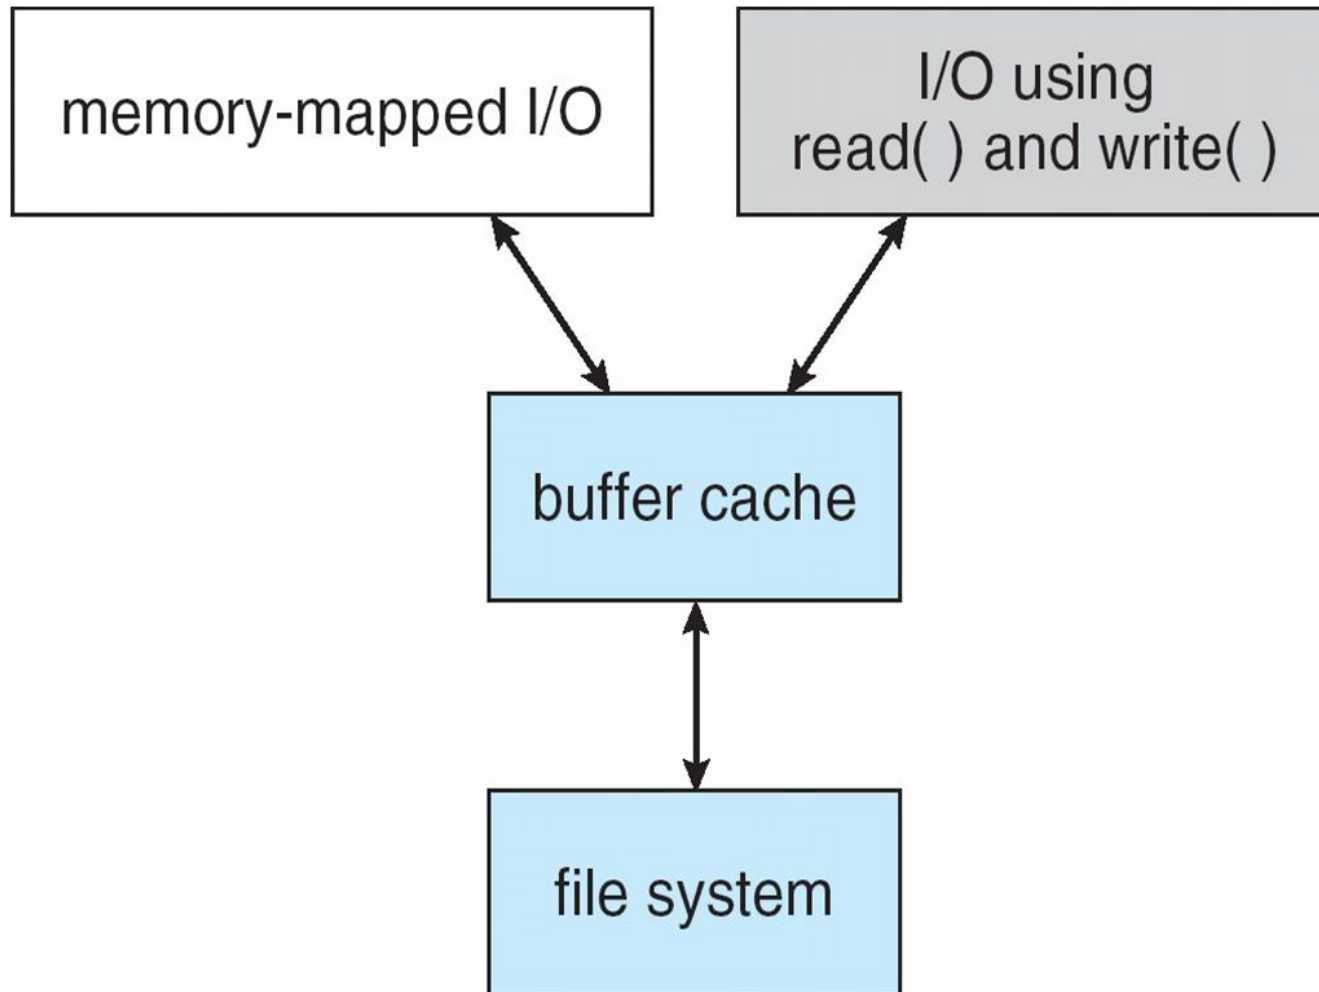
---

- A unified buffer cache uses **the same page cache** to cache both memory-mapped pages and ordinary file system I/O
- examples: Solaris, Linux, Windows NT, 2000, XP





# I/O Using a Unified Buffer Cache







# File System Consistency

---

- Multi-step updates cause problems for consistency
  - if crash happens in the middle
- E.g. transfer \$100 from my account to Janet's
  - 1. deduct \$100 from my account
  - 2. add \$100 to Janet's account
- What happens if you crash between step 1 and 2?





# File System Consistency

- E.g. Create a file
  - 1. allocate an I-NODE, and write it to disk
  - 2. write address of the I-NODE and the file name to file directory
- If crash between the 2 steps?
  - 文件在磁盘上, 但文件夹中却没有出现
  - 所谓的” 孤儿” 文件
- 如果将两步颠倒, 则可能会出现” 魅影” 文件





# File System Consistency

- **Consistency checking** – compares data in directory structure with data blocks on disk, and tries to fix inconsistencies
  - Consistency checker—a system program running at reboot time
  - UNIX: fsck
  - MS-DOS: chkdsk
- Use system programs to **back up** data from disk to another storage device (magnetic tape, other magnetic disk, optical)
- Recover lost file or disk by **restoring** data from backup





# Log Structured File Systems

- Log structured (or journaling) file systems record each update to the file system as a transaction(事务)
- All transactions are written to a log
  - A transaction is considered committed once it is written to the log
  - However, the file system may not yet be updated
- The transactions in the log are asynchronously written to the file system
  - When the file system is modified, the transaction is removed from the log
- If the file system crashes, all remaining transactions in the log must still be performed



# End of Chapter 11

---

